Hamilton cycles in pseudorandom graphs

Stefan Glock* David Munhá Correia[†] Benny Sudakov[†]

Abstract

Finding general conditions which ensure that a graph is Hamiltonian is a central topic in graph theory. An old and well known conjecture in the area states that any d-regular n-vertex graph G whose second largest eigenvalue in absolute value $\lambda(G)$ is at most d/C, for some universal constant C>0, has a Hamilton cycle. In this paper, we obtain two main results which make substantial progress towards this problem. Firstly, we settle this conjecture in full when the degree d is at least a small power of n. Secondly, in the general case we show that $\lambda(G) \leq d/C(\log n)^{1/3}$ implies the existence of a Hamilton cycle, improving the 20-year old bound of $d/\log^{1-o(1)} n$ of Krivelevich and Sudakov. We use in a novel way a variety of methods, such as a robust Pósa rotation-extension technique, the Friedman-Pippenger tree embedding with rollbacks and the absorbing method, combined with additional tools and ideas.

Our results have several interesting applications, giving best bounds on the number of generators which guarantee the Hamiltonicity of random Cayley graphs, which is an important partial case of the well known Hamiltonicity conjecture of Lovász. They can also be used to improve a result of Alon and Bourgain on additive patterns in multiplicative subgroups.

1 Introduction

A Hamilton cycle in a graph G is a cycle passing through all the vertices of G. If it exists, then G is called Hamiltonian. Being one of the most central notions in Graph Theory, it has been extensively studied by numerous researchers, see e.g., [1, 13, 17, 19, 20, 25, 29, 34, 37, 38, 40, 44], and the surveys [28, 39]. In particular, the problem of deciding Hamiltonicity of a graph is known to be NP-complete and thus, finding general conditions which ensure that G has a Hamilton cycle is one of the most popular topics in Graph Theory. For instance, two famous theorems of this nature are the celebrated result of Dirac [23], which states that if the minimum degree of an n-vertex graph G is at least n/2, then G contains a Hamilton cycle, and the criterion of Chvátal and Erdős [17] that a graph is Hamiltonian if its connectivity number is at least as large as its independence number.

In fact, most of the classical criteria for Hamiltonicity focus on rather dense graphs. A prime example of this is clearly Dirac's theorem stated above, but also the Chvátal-Erdős condition requires the graph to be relatively dense, of average degree $\Omega(\sqrt{n})$. In contrast, sufficient conditions that ensure Hamiltonicity of sparse graphs seem much more difficult to obtain. A natural starting point towards this topic is to consider sparse random graphs, to which a lot of research has been dedicated in the last 50 years. In a breakthrough paper in 1976, Pósa [44] proved that the binomial random graph model G(n,p) with $p \geq C \log n/n$ for some large constant C almost surely contains a Hamilton cycle. In doing so, he invented the influential rotation-extension technique for finding long cycles and paths, which has found numerous further applications since then. Pósa's result was later refined by Korshunov [32] and in 1983, a more precise threshold for Hamiltonicity was obtained by Bollobás [12] and Komlós and Szemerédi [31], who independently showed that if $p = (\log n + \log \log n + \omega(1))/n$, then G(n,p) is almost surely Hamiltonian. It is a standard exercise to note that this is essentially tight - indeed, if $p = (\log n + \log \log n - \omega(1))/n$, then G(n,p) almost surely has a vertex with degree at most 1, and hence is not Hamiltonian. In parallel,

^{*}Fakultät für Informatik und Mathematik, Universität Passau, Germany. Email: stefan.glock@uni-passau.de.

[†]Department of Mathematics, ETH, Zürich, Switzerland. Research supported in part by SNSF grant 200021_196965. Emails: {david.munhacanascorreia, benjamin.sudakov}@math.ethz.ch.

significant attention has also been given to the Hamiltonicity of the random d-regular graph model $G_{n,d}$ - it is known that $G_{n,d}$ almost surely contains a Hamilton cycle for all values of $3 \le d \le n - 1$. For this result, the reader is referred to Cooper, Frieze and Reed [18] and Krivelevich, Sudakov, Vu and Wormald [36] and their references.

Given the success of the study of Hamilton cycles in sparse random graphs, it became natural to then consider pseudorandom graphs, which are deterministic graphs that resemble random graphs in various important properties. A convenient way to express pseudorandomness is via spectral techniques and was introduced by Alon. An (n, d, λ) -graph is an n-vertex d-regular graph G whose second largest eigenvalue in absolute value, $\lambda(G)$, is such that $\lambda(G) \leq \lambda$. Roughly speaking, $\lambda(G)$ is a measure of how "smooth" the edge-distribution of G is, and the smaller its value, the closer to "random" G behaves. The reader is referred to [35] for a detailed survey concerning pseudorandom graphs.

In a rather influential paper, Krivelevich and Sudakov [33] employed Pósa's rotation-extension technique to prove the very general result that (n, d, λ) -graphs are Hamiltonian, provided λ is significantly smaller than d. Precisely, they showed that if n is sufficiently large, then

$$d/\lambda \ge \frac{1000 \log n(\log \log \log n)}{(\log \log n)^2} \tag{1}$$

guarantees that any (n, d, λ) -graph contains a Hamilton cycle.

This result has found numerous applications in the last 20 years towards some well-known problems, some of which we will discuss later. Given its significance and generality, it leads to the very natural and fundamental question of whether a smaller ratio of d/λ is already sufficient to imply Hamiltonicity. Krivelevich and Sudakov [33] conjectured that it should suffice that d/λ is only a large enough constant.

Conjecture 1.1. There exists an absolute constant C > 0 such that any (n, d, λ) -graph with $d/\lambda \geq C$ contains a Hamilton cycle.

Despite the plethora of incentives, there has been no improvement until now on the Krivelevich and Sudakov bound stated in (1). In this paper, we make significant progress towards Conjecture 1.1 in two ways. First, we improve on the Krivelevich and Sudakov bound in general by showing that a spectral ratio of order $(\log n)^{1/3}$ already guarantees Hamiltonicity.

Theorem 1.2. There exists a constant C > 0 such that any (n, d, λ) -graph with $d/\lambda \ge C(\log n)^{1/3}$ contains a Hamilton cycle.

The proof of the above result will rely on the Pósa rotation-extension method with various new ideas. Namely, we will need to develop some techniques in order to use this method in a *robust* manner. The reader is referred to Section 2 for an outline of the proof.

Secondly, we confirm Conjecture 1.1 in full when the degree is polynomial in the order of the graph.

Theorem 1.3. For every constant $\alpha > 0$, there exists a constant C > 0 such that any (n, d, λ) -graph with $d \geq n^{\alpha}$ and $d/\lambda \geq C$ contains a Hamilton cycle.

In fact, Theorem 1.3 is a corollary of a more general statement that we will prove (Theorem 7.3). This in particular states that (n, d, λ) -graphs with linearly many vertex-disjoint cycles are Hamiltonian.

Before discussing the applications of both of our main results, we refer the reader to Section 2 for an outline of the proof of these, as well as, of the structure of the paper.

1.1 Applications and related problems

Both Theorem 1.2 and Theorem 1.3 immediately yield improvements in several applications which made use of the result of Krivelevich and Sudakov. These will be discussed in a detailed manner in Section 8. One

application is an important special case of a famous open question of Lovász [41] from 1969 concerning the Hamiltonicity of a certain class of well-behaved graphs (see e.g., [21] and its references for more background on the problem).

Conjecture 1.4. Every connected vertex-transitive graph contains a Hamilton path, and, except for five known examples, a Hamilton cycle.

Since Cayley graphs are vertex-transitive and none of the five known exceptions in Lovász's conjecture is a Cayley graph, the conjecture in particular includes the following, which was asked much earlier in 1959 by Rapaport Strasser [45].

Conjecture 1.5. Every connected Cayley graph is Hamiltonian.

For these conjectures, a proof is currently out of sight. Indeed, notable progress towards them in their full generality are a result of Babai [8] that every vertex-transitive n-vertex graph contains a cycle of length $\Omega(\sqrt{n})$ (see [22] for a recent improvement) and a result of Christofides, Hladký and Máthé [16] that every vertex-transitive graph of linear minimum degree contains a Hamilton cycle.

Given this, it is natural to consider the "random" version of Conjecture 1.5. Indeed, Alon and Roichman [6] showed that in any group G, a random set S of $O(\log |G|)$ elements is such that the Cayley graph generated by them, $\Gamma(G, S)$, is almost surely connected. Therefore, a particular instance of Conjecture 1.5 is to show that $\Gamma(G, S)$ is almost surely Hamiltonian, which is itself a conjecture of Pak and Radoičić [43]. In fact, this relates directly to Conjecture 1.1 since it can be shown, generalizing the result of Alon and Roichman, that if $|S| \geq C \log |G|$ for some large constant C, then $\Gamma(G, S)$ is almost surely an (n, d, λ) -graph with $d/\lambda \geq K$ for some large constant K. Hence, Conjecture 1.1 would imply the Hamiltonicity of $\Gamma(G, S)$. In Section 8.1, improving on several earlier results [15, 33, 42] we will show how Theorem 1.2 can be used to prove that if |S| is of order $\log^{5/3} n$, then $\Gamma(G, S)$ is almost surely Hamiltonian (see Theorem 8.2). In the same section, we will also give an improved bound on a related problem of Akbari, Etesami, Mahini, and Mahmoody [4] concerning Hamilton cycles in coloured complete graphs which use only few colours.

Another application of our results concerns one of the central themes in Additive Combinatorics, the interplay between the two operations sum and product. A well-known fact in this area is that any multiplicative subgroup A of the finite field \mathbb{F}_q of size at least $q^{3/4}$ must contain two elements x, y such that x + y also belongs to A. Motivated by this, Alon and Bourgain [4] studied more complex additive structures in multiplicative subgroups. In particular, they proved that when a subgroup has size $|A| \geq q^{3/4}(\log q)^{1/2-o(1)}$, then there is a cyclic ordering of the elements of A such that the sum of any two consecutive elements is also in A. Using Theorem 1.3, we can improve on Alon and Bourgain's result, showing that the additional polylog-factor can be avoided. This shows that when |A| is of order $q^{3/4}$, not only does it contain $x, y, x + y \in A$ but also much more complex structures.

Finally, in Section 8.3, we give an application of our techniques to another problem related to Conjecture 1.5. Motivated by this conjecture, Pak and Radoičić [43] showed that every group G has a set of generators S of size at most $\log_2 |G|$ such that the Cayley graph $\Gamma(G, S)$ is Hamiltonian, which is optimal since there are groups that do not have generating sets of size smaller than $\log_2 |G|$. Since their proof relies on the classification of finite simple groups, they asked to find a classification-free proof of this result. Using the methods developed for the proof of Theorem 1.3 we give a classification-free proof that there is always such a set S with $|S| = O(\log n)$.

2 Outline

The proofs of Theorems 1.2 and 1.3 are quite different and in this section we give an outline of the main ideas in these proofs.

2.1 Robust Pósa rotation and Theorem 1.2

As mentioned in the introduction, the proof of Theorem 1.2 relies on Pósa's rotation-extension technique. This is a method used for finding Hamilton cycles in connected graphs. The general argument takes a longest path P in an expanding graph G and uses the rotation-extension technique to close the path P into a cycle, that is, it finds a cycle C on the same vertex set V(P). If |P| = |G|, then C is a Hamilton cycle, as desired. Otherwise, since G is connected there is a vertex $v \in V(G) \setminus C$ and an edge from v to C; this would give a longer path P' than P, contradicting the maximality of P. Clearly the crucial step of the argument above is in closing the path P into a cycle which is where the rotation-extension technique comes into play. Let us then briefly describe this beautiful method introduced by Pósa (a much more detailed take on it is made in Section 4.1).

Informally, given a path $P = (v_1, v_2, ..., v_l)$, an edge $v_i v_l$ allows one to construct a new path $P' = (v_1, v_2, ..., v_i, v_l, v_{l-1}, ..., v_{i+1})$ which also has v_1 as an endpoint. This is called a rotation of P with fixed endpoint v_1 and note that it "creates a new endpoint v_{i+1} " in that P' is a $v_1 v_{i+1}$ -path also with vertex set V(P). We can then continue this procedure, rotating P' also with fixed endpoint v_1 to create a new path P'' which has a new endpoint v_j , and so on. If a path Q is the result of t consecutive such rotations, we will refer to it as a t-rotation of P with fixed endpoint v_1 . Notice also that if P is a longest path in G, then each one of these new endpoints must have all their neighbours in V(P) - otherwise, a neighbour outside of V(P) would create a path longer than P. Using this fact we can then create many new paths and endpoints. Indeed, since our graph G has nice expansion properties we can, from starting at a longest path P and fixing one of its endpoints v_1 , find O(n) vertices v_j such that there is a v_1v_j -path also on the vertex set V(P) which is a t-rotation of P for some $t = O(\log n)$ (for a precise statement of this, see Lemma 4.6 and Corollary 4.7).

Our goal is now to close P into a cycle, that is, find a spanning cycle on the vertex set V(P). A reasonable first approach is as follows. By the previous paragraph, the rotation technique allows to find a set X of size $\Omega(n)$ such that every vertex $v_i \in X$ is such that there is a v_1v_i -path P_i with vertex set V(P). Similarly, for each $v_j \in X$, we can find a set Y_j of size $\Omega(n)$ such that every vertex $v_k \in Y_j$ is such that there is a $v_k v_i$ -path $P_{k,j}$ with vertex set V(P). Clearly, the set Y_i is dependent on the vertex $v_i \in X$ which is chosen. Nevertheless, it might be that there exists a large set $X' \subseteq X$ such that the intersection of all sets Y_j with $v_j \in X'$ is also large. Then, denoting this intersection by Y', it will be the case that for each $a \in X', b \in Y'$ there is an ab-path with vertex set V(P). Using that G is a pseudorandom graph, we can then find an edge ab with $a \in X'$ and $b \in Y'$ - this creates a cycle as desired. For example, one can observe that such X', Y' can be found if we can: rotate consecutively with fixed endpoint v_1 and only using endpoints in $\{v_{l/2+1}, \ldots, v_{l-1}, v_l\}$, that is, only using edges induced by these vertices; and then rotate consecutively with fixed endpoint v_l (or any endpoint in $\{v_{l/2+1}, \ldots, v_{l-1}, v_l\}$) and only using endpoints in $\{v_1, \ldots, v_{l/2}\}$. Indeed, then since the two sets of rotations do not interfere with each other, we will get two large sets $X' \subseteq \{v_{l/2+1}, \dots, v_{l-1}, v_l\}, Y' \subseteq \{v_1, \dots, v_{l/2}\}$ such that for each $a \in X', b \in Y'$ there is a desired ab-path. However, it is clearly not the case in general that such a situation will hold. This is the main issue that we have to resolve. Nevertheless, the objective is to, using rotations, find disjoint sets X, Y which are large enough, such that for every $a \in X, b \in Y$, there is an ab-path with vertex set V(P). Then, by pseudorandomness there is an edge between X, Y which will create the desired cycle.

As we saw with the example in the previous paragraph, in order to do this, it will be convenient to control the way we rotate. With this in mind, we will first clean the path P (this is done in Lemma 5.2). Informally, if G is an (n, d, λ) -graph and P is large, we can find a collection of disjoint intervals in the path which cover almost all the vertices of P such that almost every vertex v in each of these intervals has many neighbours in several other intervals. Clearly, this will allow us to rotate in various different manners, avoiding certain situations. Indeed, let us informally describe how the proof goes after this cleaning is done. Suppose for simplicity that P is fully partitioned into k intervals of size |P|/k and

that all the vertices have many neighbours in at least εk intervals, for some appropriate ε . We now fix a large sub-path A at the start of the path containing the vertex v_1 which contains roughly $\varepsilon k/2$ intervals. Observe already that since A has that size, every vertex has lots of neighbours outside of A (it has many neighbours in at least $\varepsilon k/2$ intervals which are disjoint to A). This implies that we can always rotate from any vertex while avoiding A - let us denote this property by \mathbf{P} . Also, using the pseudorandomness of G we can find a subset $A' \subseteq A$ of size 0.99|A| which has large minimum degree and has the property that every vertex in A' has lots of neighbours such that both of their neighbours in P also belong to A'. Note that this implies that if we are doing rotations and our current endpoint belongs to A', then we can continue rotating while "staying" in $A' \subseteq A$, that is, only creating new endpoints which belong to A' - let us call this property \mathbf{Q} .

Now, using property \mathbf{P} , we can first rotate with fixed endpoint v_l while avoiding A to find many new endpoints. Then, the pseudorandomness of G guarantees that we can find an edge between one of these endpoints and A', giving a new endpoint $z \in A'$. Moreover, crucially note that since we only "touched" A at this last step, we now have a zv_l -path P' such that A is broken into only two disjoint sub-paths of P' (one at the start containing z and one at the middle - see Figure 1 where these paths are depicted in red) and thus, in particular, A' still has property Q. Let us denote the interval of P' between these two sub-paths of A as I. Using property \mathbf{Q} we can then start rotating P' with fixed endpoint v_l while staying in $A' \subseteq A$. Let the set of endpoints created as such be denoted as X, so that for each $x \in X$ we created an xv_l -path P_x . Since $P' \setminus A$ was avoided in these rotations, observe that the only difference between $P_x \setminus A$ and $P_{x'} \setminus A$ (for two different $x, x' \in X$) in the paths $P_x, P_{x'}$ is the direction in which the interval I is traversed in the orientation $x, x' \to v_l$ (see Figure 1 for an illustration). Then, by possibly reducing the size of X by half, we can assume that this direction is the same for all endpoints in X - let us denote this property by **R**. Finally, we can now, for some $x \in X$, use property **P** to rotate with fixed endpoint x only using $P' \setminus A$. Then, since these rotations avoid A, property R implies that every endpoint y this creates is also an endpoint for all other $x' \in X$, that is, for all $x' \in X$ there is an x'y-path on the vertex set V(P). Letting Y denote this set of endpoints y, if X, Y are large enough then the pseudorandomness of G implies that there is an edge xy with $x \in X, y \in Y$, thus creating a cycle as desired.

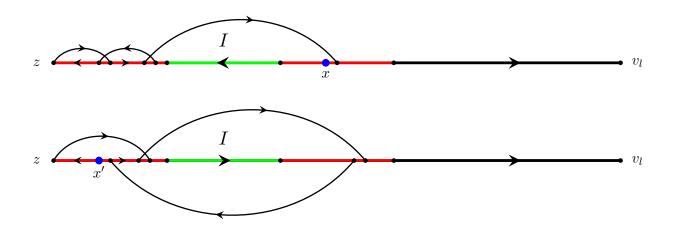


Figure 1: The zv_l -path P' is depicted, the set A is pictured in red and the interval I in green. Two rotations with fixed endpoint v_l which "stay" in A (using property \mathbf{Q}) are depicted with the interval I being traversed in different directions in the resulting paths $P_x, P_{x'}$ which are depicted using arrows in the direction $x, x' \to v_l$.

Clearly, the previous discussion is only informal and lacks a lot of important details that we have to deal with in order to do the proof. Namely, at the start it will not be the case that we can partition all of P

into intervals so that every vertex has many neighbours in several intervals (and so, in particular, we will not have property \mathbf{P}). In particular, the endpoints of the path P will not necessarily have neighbours in many intervals. Therefore, we will first have to perform some rotations in order to get endpoints which do satisfy this. This will however imply that some changes to the path will be made, which requires that when we go through the arguments layed out before, we will need to rotate while taking these changes into account. This is an instance of where Lemma 4.8 comes in handy, which is the main technical lemma allowing us to rotate in a *robust manner*. In Section 4 we discuss the standard Pósa rotation technique, state and prove our technical lemma as well as discuss our method for closing paths into cycles. Then, in Section 5, we give the cleaning lemma discussed before (Lemma 5.2) and prove Theorem 1.2.

2.2 Graphs with many cycles and Theorem 1.3

We now discuss the proof of Theorem 1.3 along with the more general statement (Theorem 7.3) that every (n, d, λ) -graph with many disjoint cycles contains a Hamilton cycle. Indeed, if $d \geq n^{\alpha}$ for some constant $\alpha > 0$, then it is not difficult to show that the graph contains a collection of εn many vertex-disjoint cycles, for some constant ε dependent on α .

For simplicity, let us assume here that all of these cycles are triangles and denote them as $a_ib_ic_i$ for $1 \le i \le \varepsilon n$. First we will find a path P_1 which connects linearly many of the triangles, that is, for each such triangle it contains the edge a_ic_i but not the vertex b_i - this can be achieved using rather standard arguments, namely a directed version of the Depth First Search algorithm (see Lemma 3.17). We have now build a flexible set B of linear size consisting of all the vertices b_i , meaning that we can flexibly use the vertices of B to build structure in the remainder of the graph, and be assured that afterwards, the vertices b_i which we have not used can be "absorbed" into the path P_1 by replacing a_ic_i with the path $a_ib_ic_i$. Moreover, by discarding a few vertices, we can find a set $B' \subseteq B$, still of linear size, such that G[B'] has large minimum degree, which will imply that it inherits some expansion properties from G, and also such that the endpoints of P_1 have large degree into B'.

Let Y be the set of vertices not contained in P_1 and B'. With some additional care, we can ensure that G[Y] has large minimum degree and therefore it has nice expansion properties. The goal is now to partition G[Y] into few paths, more precisely, at most $o(n/\log n)$ vertex-disjoint paths which cover all vertices of Y, such that the endpoints of all of these paths have neighbours in the flexible set - this is all done in Section 6. For simplicity, let us start describing briefly how the first property is achieved and then explain how these two properties can be used to find a Hamilton cycle. At the end, we discuss the second property.

One standard argument which has been used in the literature for finding linear forests with few paths in these well-behaved graphs is to find a 2-factor with few cycles by showing that the total number of 2-factors is larger than the number of 2-factors with many cycles using estimates on the permanent as well as facts about (pseudo-)random graphs. Unfortunately this is not applicable to very sparse graphs and instead we need to come up with a different argument. Interestingly, we find a convenient shortcut to this problem. We will use the expanding properties of G[Y] to find in it a spanning regular subgraph and then invoke a result of Alon [3] on the linear arboricity conjecture to find the desired collection of paths. Now, let P_2, \ldots, P_r be these paths which cover Y, with $r = o(n/\log n)$, and let u_i, v_i denote the endpoints of P_i - assume that each one of these endpoints has neighbours in the flexible set B'. We want to now weave the paths together with the path P_1 into a Hamilton cycle. More precisely, for each i, we want to find a path which connects v_i to u_{i+1} (indices modulo r), whose inner vertices lie in B', such that all these paths are vertex-disjoint. Observe that this would indeed yield the desired Hamilton cycle. Since B' enjoys nice expansion properties, we can generally assume to find paths of size $\log n$ between any given pair of vertices. Since we only need to connect $o(n/\log n)$ pairs, in principal there would be enough space in B' to find all the connections vertex-disjointly. This vertex-disjoint paths problem has been extensively studied in the literature. We benefit here tremendously from a recent "roll-back technique" which allows one to indeed find all these connections in expanding graphs (see [24]). Roughly speaking, this technique allows us, for a given pair of vertices we want to connect, to build two binary trees rooted at both vertices in a controlled fashion (using a concept of Friedman and Pippenger [26]); then, when these trees are large enough, we are guaranteed by the pseudorandomness of G to find an edge between them which then obviously leads to a path connecting the pair. The astonishing thing is that we can now roll back, which means that we can demolish the two trees until only the desired path is left, while having a guarantee that the resulting embedding is still good in the sense that we can simply start connecting the next pair disjointly in the same way. In fact, we will have to use this technique with an additional twist. Namely, in general we cannot find triangles at the start of the proof, but cycles which might be longer. For example, say that we have cycles a_i, b_i, c_i, d_i of length four. Now the path P_1 will connect some of these in the sense that P_1 contains the edges a_id_i , but does not contain the vertices b_i, c_i . We still have a flexible set now, but we have to use it in such a way that for each i, either both b_i, c_i are absorbed (and in a particular way) or neither. In Section 3.5, we deal with this additional complication.

Finally, let us briefly discuss the second property that the endpoints of P_2, \ldots, P_r have neighbours in the flexible set B' - indeed, otherwise, we could not connect them using this set. For this, let X be the subset of Y consisting of those vertices which do not have enough neighbours in the flexible set - since G is pseudorandom, we are guaranteed that X is not large. When finding the paths P_2, \ldots, P_r , we have to make sure that their endpoints avoid X. In order to achieve this, we will first "absorb" the whole of X into $Y \setminus X$. This uses a novel idea (see Lemma 6.4) on how to cover such a bad set X which might be very useful for future applications (in particular since we do not need the polynomial degree assumption here). Precisely, we can find a collection of disjoint paths which cover X and whose endpoints are in $Y \setminus X$. Then, we will like before use the result of Alon [3] to extend this collection into a linear forest with few paths whose endpoints are in $Y \setminus X$, as desired. Clearly, doing this when $X \neq \emptyset$ requires various new ideas.

The part of the paper devoted to Theorem 1.3 is organized as follows. In Section 3.5, we present all the connecting tools we require - both to find at the start the path P_1 connecting the cycles as well as to connect the paths P_2, \ldots, P_r at the end using the flexible set to construct the Hamilton cycle. Subsequently, in Section 6 we discuss finding the desired linear forest in our context with few paths with good endpoints. Then, in Section 7 we prove the main theorem.

3 Preliminaries

3.1 Notation

We mostly use standard graph theoretic notation. Let G be a finite graph. Denote by V(G) its vertex set, and let $S_1, S_2 \subseteq V(G)$. We denote by $G[S_1]$ the subgraph of G induced by S_1 , and by $E[S_1, S_2]$ the set of edges with one endpoint in S_1 and the other in S_2 . For a vertex x in a graph G, we let $N_G(x)$ denote the neighbourhood of x in G. For a set S of vertices in a graph G, we let $\Gamma_G(S)$ (or just $\Gamma(S)$ if G is clear from the context) denote the neighbourhood of X, that is $\Gamma_G(X) := \bigcup_{x \in X} N_G(x)$, and we denote by $N_G(X)$ the external neighbourhood of X, that is $N_G(X) = \Gamma_G(X) \setminus X$. If G is a directed graph, then we can also define $N_G^+(S), \Gamma_G^+(S), N_G^-(S), \Gamma_G^-(S)$ to be the analogous sets considering out-or in-neighbourhoods. Usually n denotes the number of vertices of a given graph, and we assume in all our statements that n is sufficiently large. An event holds with high probability if its probability tends to 1 as $n \to \infty$. Given a collection \mathcal{M} of vertex disjoint sets $S \subseteq V(G)$ in a graph G, we will use $V(\mathcal{M})$ to denote the set of all vertices in such sets, that is, $V(\mathcal{M}) := \bigcup_{S \in \mathcal{M}} S$.

3.2 Some definitions and standard tools

We now give two useful definitions concerning pseudorandom graphs.

Definition 3.1. Let G be an (n, d, λ) -graph. A subset $S \subseteq V(G)$ is said to be δ -clean if G[S] has minimum degree at least δ . In the case that $\delta \geq d|S|/4n$, we say simply that S is clean.

Definition 3.2. A graph G is said to be an (s, K)-expander (or (s, K)-expanding) if every subset $S \subseteq V(G)$ of size at most s is such that $|N_G(S)| \ge K|S|$. Similarly, a directed graph G is (s, K)-out-expanding if for every subset $S \subseteq V(G)$ of size $|S| \le s$ we have $|N_G^+(S)| \ge K|S|$.

Now we will state various well-known results that we will use throughout the paper. The first two are standard probabilistic tools.

Lemma 3.3 (Chernoff's bound - see e.g., [7]). Let X be the sum of independent random variables X_1, \ldots, X_n such that $0 \le X_i \le k$ for each i. Then, for all $0 < \varepsilon < 1$,

$$\mathbb{P}\left(X < (1 - \varepsilon)\mathbb{E}[X]\right), \mathbb{P}\left(X > (1 + \varepsilon)\mathbb{E}[X]\right) \le e^{-\varepsilon^2\mathbb{E}[X]/3k^2}.$$

Lemma 3.4 (Lovász local lemma - see e.g., [7]). Let A_1, \ldots, A_n be events in a probability space and suppose that each event A_i is mutually independent of a set of all the other events A_j but at most d. If $\mathbb{P}(A_i) \leq p$ for all $i \in [n]$ and $ep(d+1) \leq 1$, then

$$\mathbb{P}\left(\bigcap_{i}\overline{A_{i}}\right) > \left(1 - \frac{1}{d+1}\right)^{n}.$$

Next we state the classical result of Hall concerning matchings in bipartite graphs.

Theorem 3.5 (Hall's theorem). Let G be a bipartite graph with bipartition $V(G) = A \cup B$ and such that

$$|N(S)| \ge |S| \quad \forall S \subseteq A.$$

Then G has a matching covering A.

We also state a simple folklore result to find short cycles in sufficiently dense graphs.

Theorem 3.6. Every n-vertex graph with at least $n^{1+1/l}$ edges contains a cycle of length at most 2l.

Finally, we state another folklore result on the spectral gap of (n, d, λ) -graphs.

Theorem 3.7. Every
$$(n, d, \lambda)$$
-graph must have $\lambda \geq \sqrt{d \cdot \frac{n-d}{n-1}}$.

We remark that by Dirac's theorem stated in the introduction, if $d \ge n/2$ then the graph is Hamiltonian. Hence, we can from now on assume that d < n/2 and thus, the above theorem tells us that $\lambda \ge \sqrt{d/2}$.

3.3 Expansion properties of pseudorandom graphs

Here we state and prove several facts about (n, d, λ) -graphs which relate to their expansion properties. As it is usual, the starting point is the celebrated expander mixing lemma, introduced by Alon and Chung [5], which has since then found numerous applications (the interested reader is referred to Chapter 9 of the book of Alon and Spencer [7]). Below, we state the lemma together with three well-known corollaries. For a proof of these, the reader can consult Section 2 of Krivelevich and Sudakov [33].

Lemma 3.8 (Expander mixing lemma). Let G be an (n, d, λ) -graph. Then,

(1) For every two subsets $A, B \subseteq V(G)$, the number e(A, B) of edges with one endpoint in A and the other in B satisfies

$$\left| e(A,B) - \frac{d|A||B|}{n} \right| \le \lambda \sqrt{|A||B|}.$$

(2) For every set $A \subseteq V(G)$,

$$\left| e(A) - \frac{d|A|^2}{2n} \right| \le \frac{\lambda|A|}{2}.$$

Moreover, the following three properties hold.

- (3) Every $A \subseteq V(G)$ of size at most $\lambda n/d$ has $e(A) \leq \lambda |A|$.
- (4) For every $A, B \subseteq V(G)$ such that $|A||B| > (\lambda n/d)^2$, there exists an edge with one endpoint in A and the other in B.
- (5) If $d > 2\lambda$, then G is connected.

We now give the two crucial lemmas concerning expansion in pseudorandom graphs. These will be corollaries of the expander mixing lemma above. The first states, informally, that as long as there is large enough minimum degree (of order at least λ), we have good vertex-expansion.

Lemma 3.9. Let G be an (n, d, λ) -graph, $G' \subseteq G$ a subgraph and $10\lambda \le \delta \le d$. Suppose that there are subsets of vertices A, B with $|A| \le \delta n/100d$ and such that every vertex $v \in A$ has $|N_{G'}(v) \cap B| \ge \delta$. Then,

$$|N_{G'}(A) \cap B| \ge \min\left(\delta^2 |A|/8\lambda^2, \delta n/10d\right) \ge 10|A|.$$

Proof. Let $T := N_{G'}(A) \cap B$. Firstly, since $|A| \le \delta n/100d$ note that by part (2) of Lemma 3.8 we have $e_G(A) \le (\delta/200 + \lambda/2)|A| \le 0.15\delta|A|$. In turn, by assumption and since $\delta \ge 10\lambda$, this implies that

$$e_G(A, T) \ge e_{G'}(A, T) \ge \delta |A| - 2e_G(A) \ge 0.7\delta |A|.$$

Now, suppose first that $|A| \leq \lambda^2 n/\delta d$ and for sake of contradiction, also that $|T| < \delta^2 |A|/8\lambda^2$. Then, part (1) of Lemma 3.8 implies that

$$e_G(A,T) \leq \frac{d|A||T|}{n} + \lambda \sqrt{|A||T|} < \left(\frac{\delta^2|A|d}{8\lambda^2 n} + \lambda \cdot \frac{\delta}{2\lambda}\right)|A| \leq \left(\frac{\delta}{8} + \frac{\delta}{2}\right)|A| < 0.7\delta|A|,$$

where the third inequality follows since $|A| \leq \lambda^2 n/\delta d$. This is a contradiction to the previous observation and thus, as desired, $|T| \geq \delta^2 |A|/8\lambda^2$. Finally, suppose that $|A| > \lambda^2 n/\delta d$. By considering a subset $A' \subseteq A$ of size $\lambda^2 n/\delta d$, we have by before that $|T| \geq \delta^2 |A'|/8\lambda^2 - |A| \geq \delta n/8d - |A| > \delta n/10d$, where we are using that $|A| \leq \delta n/100d$.

The following is an immediate corollary of the above.

Corollary 3.10. Let G be an (n, d, λ) -graph and $S \subseteq V(G)$ be a clean subset of size at least $1000\lambda n/d$. Then, G[S] is an (|S|/400, 10)-expander.

Next we give another lemma which is a corollary of Lemma 3.8. This will tell us that every sufficiently large set contains a large clean subset. This is important since by Corollary 3.10, it implies that it will have good expansion properties.

Lemma 3.11. Let G be an (n,d,λ) -graph and $S \subseteq V(G)$. Then, there are at most $4\lambda^2 n^2/d^2|S|$ vertices u such that $|N(u) \cap S| < d|S|/2n$. Moreover, if $|S| \ge 5\lambda n/d$, there is a clean subset $S' \subseteq S$ of size at least $|S| - \lambda n/d$.

Proof. Let U denote the set of vertices u such that $|N(u) \cap S| < d|S|/2n$. Then, $e(U,S) \le d|U||S|/2n$. At the same time, part (1) of Lemma 3.8 implies that $e(U,S) \ge d|U||S|/n - \lambda \sqrt{|U||S|}$. Therefore,

$$\lambda \sqrt{|U||S|} \ge d|U||S|/2n,$$

and so, $|U| \leq 4\lambda^2 n^2/d^2|S|$ as desired.

For the second part, do the following process. Start with $S_0 := S$; at each step $i \geq 1$, consider $G[S_{i-1}]$: if it contains some vertex u_i with degree less than d|S|/4n, then remove it, that is, define $S_i = S_{i-1} \setminus \{u_i\}$ and continue; otherwise, stop the process. Now, it is clear that if the process stops at some step $i \leq \lambda n/d$, then the desired subset $S' := S_i \subseteq S$ is produced. For sake of contradiction, suppose then that it only stops later. Consider the set S_i for $i := \lambda n/d + 1$ and define $U := \{u_1, \ldots, u_{i-1}\}$, which is of size at least $\lambda n/d$. Note that $|S_i| = |S| - \lambda n/d \geq \max(|S|/2, 4\lambda n/d)$ and further, by definition, every vertex in U must have less than $d|S|/4n < d|S_i|/2n$ neighbours in S_i . However, this contradicts the first part of the corollary since

$$|U| > \lambda n/d \ge 4\lambda^2 n^2/d^2 |S_i|$$

where the second inequality holds since, as noted before, $|S_i| \ge 4\lambda n/d$.

We remark that with the same argument one can easily generalize the second part of the above lemma in the following manner.

Lemma 3.12. Let G be an (n, d, λ) -graph and S a collection of disjoint pairs $\{x_i, y_i\} \subseteq V(G)$. If $|S| \geq 3\lambda n/d$, there is a sub-collection $S' \subseteq S$ of size at least $|S| - \lambda n/d$ such that V(S') is clean.

3.4 Directed graphs and the (directed) Friedman-Pippenger technique

This next section is devoted to stating lemmas concerning finding sparse structures such as paths and trees more generally, as well as connecting structures, in expanding directed graphs, which will often appear in our proofs as auxiliary graphs. We start with a well-known result of Ben-Eliezer, Krivelevich and Sudakov [10] which gives a sufficient condition for the existence of long paths in directed graphs. Its proof relies on the depth-first-search (DFS) algorithm.

Lemma 3.13 ([10]). Let H be an n-vertex digraph and suppose that for all disjoint sets $S, T \subseteq V(H)$ of size k there exists an edge directed from S to T. Then, H contains a directed path of length n-2k+1.

We now describe an embedding machinery in expanding directed graphs. This is a directed version of the Friedman-Pippenger technique with roll-backs which is used for embeddings in expander graphs. The technique of Friedman and Pippenger [26] can be used to embed trees in expanders vertex by vertex, while maintaining a certain invariant. Then, a simple but powerful observation allows one to remove leaves from this tree so that this invariant is still maintained. We refer the reader to [24] for more on the topic. First, we give the definition of this invariant - a good embedding.

Definition 3.14. Let H be a digraph and let $s, D \in \mathbb{N}$. Given a digraph F, we say that an embedding $\phi \colon F \hookrightarrow H$ is (s, D)-good if

$$|\Gamma_H^+(X) \setminus \phi(F)| \ge \sum_{v \in X} \left[D - \deg_F(\phi^{-1}(v)) \right] + |\phi(F) \cap X| \tag{2}$$

for every $X \subseteq V(H)$ of size $|X| \leq s$. Here we slightly abuse notation by setting $\deg_F(\emptyset) := 0$, i.e. if a vertex $v \in V(H)$ is not used by ϕ to embed F, then we set $\deg_F(\phi^{-1}(v)) = 0$.

Then, the argument of Friedman-Pippenger (see e.g., Theorem 2.3 and its proof in [24]) can be adapted straightforwardly towards directed graphs to give the following.

Theorem 3.15. Let F be a digraph with $\Delta(F) \leq D$ and v(F) < s, for some $D, s \in \mathbb{N}$. Suppose we are given a (2s-2,D+2)-out-expanding digraph H and a (2s-2,D)-good embedding $\phi \colon F \hookrightarrow G$. Then for every digraph F' with $v(F') \leq s$ and $\Delta(F') \leq D$ which can be obtained from F by successively adding a new vertex of in-degree 1 and out-degree 0, there exists a (2s-2,D)-good embedding $\phi' \colon F' \hookrightarrow G$ which extends ϕ .

Further, the roll-back technique can also be adapted to directed graphs.

Lemma 3.16. Suppose we are given digraphs G and F with $\Delta(F) \leq D$, and an (s, D)-good embedding $\phi \colon F \hookrightarrow G$, for some $s, D \in \mathbb{N}$. Then for every digraph F' obtained from F by successively removing a vertex of in-degree 1 and out-degree 0, the restriction ϕ' of ϕ to F' is also (s, D)-good.

Again, this can be proved by slightly adapting the undirected version - see e.g., Lemma 2.4 in [24].

3.5 Alternating paths in pseudorandom graphs

The purpose of this section is to consider the following setting. We are given a (pseudorandom) graph G and a collection \mathcal{M} of disjoint pairs $\{x,y\} \subseteq V(G)$, with the possibility that x=y - we let $G \cup \mathcal{M}$ denote the result of adding all possible and non-existing edges xy to G when $x \neq y$. We say that a path P in $G \cup \mathcal{M}$ is \mathcal{M} -alternating if it is of the form $e_1f_1e_2f_2\dots e_lf_l$ where the e_i 's are edges of G and the f_i 's represent either edges xy such that $\{x,y\}$ is a pair in \mathcal{M} or vertices x such that $\{x,x\} \in \mathcal{M}$; informally, without the appearance of pairs $\{x,x\}$, it is just a path in which there is an edge of \mathcal{M} among every two consecutive edges. The goal is now to extend some well-known results concerning paths in pseudorandom graphs to results about \mathcal{M} -alternating paths in such graphs. The first will be the following extension of Lemma 3.13.

Lemma 3.17. Let G be an n-vertex graph such that for every two disjoint subsets $S, T \subseteq V(G)$ of size k there is an edge between them. Then, $G \cup \mathcal{M}$ contains an \mathcal{M} -alternating path which uses all but 2k-1 pairs in \mathcal{M} .

The second result about alternating paths in pseudorandom graphs will be an extension of the vertexdisjoint paths problem. Roughly speaking, this topic deals with the following situation: we are given a graph G and pairs of vertices $\{a_i, b_i\}$; then, we have to find vertex disjoint paths P_i for each i connecting a_i to b_i . This problem has been extensively studied and for a detailed discussion on it the reader is referred, e.g., to [24]. In particular, Draganić, Krivelevich and Nenadov [24] recently used the Friedman-Pippenger technique to show that this problem can be solved in (n, d, λ) -graphs (and with a polynomial time online algorithm). In the next result, we extend this result to the 'alternating' version that we require.

Theorem 3.18. Let G be an (n,d,λ) -graph with $d > 500\lambda$ and $V(\mathcal{M})$ be 500λ -clean in G with $|\mathcal{M}| \geq 2000\lambda n/d$. Let \mathcal{P} be a collection of at most $\frac{|\mathcal{M}|}{100\log_2 n}$ disjoint pairs $\{a_i,b_i\}$ of vertices in $V(G)\setminus V(\mathcal{M})$ such that for all $v\in V(\mathcal{P})$ we have $|N(v)\cap V(\mathcal{M})|\geq 500\lambda$. Then there exist vertex-disjoint \mathcal{M} -alternating paths in $G\cup \mathcal{M}$ between every pair of vertices $\{a_i,b_i\}$.

Now, before going into the proofs of both of these results, we define a random auxiliary digraph H for this setting. Indeed, suppose we have a graph G and a collection \mathcal{M} of disjoint pairs $\{x,y\} \subseteq V(G)$, with the possibility that x=y. We define the random auxiliary digraph $H=H(G,\mathcal{M})$ as follows. For each pair $\{x_i,y_i\} \in \mathcal{M}$, colour uniformly at random one of the vertices red and the other blue - in the case that $x_i=y_i$ colour the vertex red and blue. After doing so, re-label these vertices so that x_i always refers to a red vertex and let X denote the set of these; let Y denote the set of blue vertices. Now, the vertex set of H is $V(H):=X\cup (V(G)\setminus V(\mathcal{M}))$ and the edges are defined as follows.

- 1. For any two vertices $x_i, x_j \in X$, put an edge directed from x_i to x_j if $x_i y_j$ is an edge in G.
- 2. For $v \notin V(\mathcal{M})$ and a vertex $x_i \in X$, put an edge from v to x_i if vy_i is an edge of G.

Crucially, constructing H in such a manner gives the following deterministic property.

Observation 3.19. If $P = x_{i_1} \to x_{i_2} \to \ldots \to x_{i_l}$ is a directed path in H, then $y_{i_1}x_{i_1}y_{i_2}x_{i_2}\ldots y_{i_l}x_{i_l}$ is an M-alternating path in $G \cup M$.

We can now prove the results stated above. The first one, Lemma 3.17 will not require the randomness involved in H and follows by showing that the condition of Lemma 3.13 is satisfied for H[X] and then applying the observation above.

Proof of Lemma 3.17. Let S,T be two disjoint subsets of X of size k. Consider the set $T' := \{y_i : x_i \in T\}$ which has size k, and note that by construction, $\overrightarrow{e}_H(S,T) = e_G(S,T')$. By the assumption on G, we then have that $\overrightarrow{e}_H(S,T) \neq \emptyset$ and therefore, Lemma 3.13 implies that H[X] contains a directed path of length |X| - 2k + 1. By Observation 3.19 this gives an M-alternating path as desired.

The second result is more technical and will indeed require the randomness of H. Precisely, we will first show that with positive probability H satisfies strong expansion properties. Because of this, we are then able to apply the directed Friedman-Pippenger machinery introduced in Section 3.4. Before starting, let us state another deterministic property that H satisfies.

Observation 3.20. Let $P = v \to x_{i_1} \to \ldots \to x_{i_l}$ and $Q = u \to x_{j_1} \to \ldots \to x_{j_t}$ be two disjoint directed paths in H with $u, v \notin V(\mathcal{M})$. Then, if $x_{i_l}x_{j_t}$ is an edge of G, we have that $vy_{i_1}x_{i_1} \ldots y_{i_l}x_{i_l}x_{j_t}y_{j_t} \ldots x_{j_1}y_{j_1}u$ is an \mathcal{M} -alternating path in $G \cup \mathcal{M}$.

Proof of Theorem 3.18. Let us fix $\delta := 500\lambda$. Let us also re-define $H := H[X \cup V(\mathcal{P})]$ - note than that $|V(H)| = |\mathcal{M}| + |V(\mathcal{P})| \le 2|\mathcal{M}|$. We will first show that H satisfies good expansion properties.

Claim. With positive probability, H is (|V(H)|/10, 5)-out-expanding.

Proof. We first show that with positive probability, $\delta^+(H) \geq \delta/5$ in H. Indeed, by assumption initially each vertex $v \in V(\mathcal{P}) \cup V(\mathcal{M})$ of G has at least δ neighbours in $V(\mathcal{M}) = X \cup Y$. It is easy to check that the random procedure defining H implies that the random variable $|N_H^+(v)|$ stochastically dominates the binomial random variable $\text{Bin}(\delta, 1/2)$. Indeed, note that if v has both elements of some pair $\{x, y\}$ in \mathcal{M} as neighbours in G, then one of them will be an in/out-neighbour in H, and if only one of these is a neighbour, then with probability 1/2 it will be an in/out-neighbour in H. Therefore, by Lemma 3.3 we have that $\mathbb{P}\left(|N_H^+(v)| < \delta/5 \mid v \in V(H)\right) = e^{-\Omega(\delta)}$. Moreover, we can further apply Lemma 3.4. Indeed, note that clearly since G is d-regular, the event $E_v := \{|N_H^+(v)| < \delta/5 \mid v \in V(H)\}$ will depend on at most $O(d^2)$ other events $E_{v'}$. Therefore, since $e^{-\Omega(\delta)} = o(d^{-2})$ (by $\delta \geq \sqrt{d/2}$), Lemma 3.4 implies that with positive probability, no event E_v holds and thus, $\delta^+(H) \geq \delta/5$.

We can now verify the expanding condition. Let $S \subseteq V(H)$ be a set of size at most |V(H)|/10 and consider $X' := \Gamma_H^+(S) \subseteq X$. Recall that for each $x_i \in X'$, there is a corresponding $y_i \in Y$ - let Y' denote the set of these. By the paragraph above, we have that for all $v \in S$, $|N_H^+(v) \cap X'| \ge \delta/5$. In turn, by the construction of H, it must be that for every edge $v \to x_i$ in H with $x_i \in X'$, there is a corresponding edge vy_i of G. Therefore, also $|N_G(v) \cap Y'| \ge \delta/5$. By Lemma 3.9, we must then have that if $|S| \le \delta n/500d$, then $|X'| = |Y'| \ge 10|S|$, as desired. When $\lambda n/d = \delta n/500d < |S| \le |V(H)|/10$, part (4) of Lemma 3.8 applied to G implies that there are at most $\lambda n/d$ vertices of Y which are not adjacent to some vertex of S - that is, $|X'| = |Y'| \ge |Y| - \lambda n/d \ge |V(H)| - \lambda n/d - |V(P)| \ge 0.9|V(H)| \ge 6|S|$, as desired. \square

Now that we have proven that H is out-expanding, we can apply the Friedman-Pippenger machinery to find the desired \mathcal{M} -alternating paths in $G \cup V(\mathcal{M})$ connecting the pairs in \mathcal{P} . Recalling Definition 3.14, we will from now on refer to a (|V(H)|/10,3)-good embedding as only a good embedding. Similarly, we will say that a directed graph is expanding if it is (|V(H)|/10,5)-out-expanding.

First note that the set $V(\mathcal{P})$ is a good embedding of an independent set in H, since H is expanding and every vertex $v \in V(\mathcal{P})$ has only out-neighbours in X, which is disjoint to $V(\mathcal{P})$. We then start by linking the first pair $\{a_1, b_1\} \in \mathcal{P}$. For this, we use Theorem 3.15 in order to find two disjoint out-oriented binary trees T_1, T_2 in H of size at least $|V(H)|/50 \ge |\mathcal{M}|/50 \ge \lambda n/d$ rooted at a_1 and b_1 respectively. Furthermore, Theorem 3.15 guarantees that these trees together with the other vertices in $V(\mathcal{P})$ form

a good embedding of an independent set of size $2|\mathcal{P}| - 2$ together with two disjoint out-oriented binary trees. Now, since $|T_1|, |T_2| \geq \lambda n/d$, part (4) of Lemma 3.8 implies that there is an edge e = uw in G[X] connecting a vertex $u \in V(T_1)$ and a vertex $w \in V(T_2)$. By Observation 3.20, this implies the existence of an \mathcal{M} -alternating path in $G \cup \mathcal{M}$ with endpoints a_1, b_1 , which we denote by P_1 . Note that P_1 uses at most $2 \log_2 n$ vertices of H.

We can now roll back in the following sense. Trivially we can remove vertices from each tree T_1, T_2 in a leaf-by-leaf manner so that at the end we are only left with the vertices of H which correspond to the path P_1 (which consist of a directed a_1u -path and a directed b_1w -path). By Lemma 3.16, we are guaranteed that the current forest remains a good embedding, now of an independent set of size 2|P|-2 and two disjoint directed paths. We then continue with the same procedure to link the second pair $\{a_2, b_2\} \in \mathcal{P}$ with an \mathcal{M} -alternating path P_2 disjoint to P_1 . Again, we first find two complete out-oriented binary trees in H rooted at a_2 and b_2 (disjoint from P_1), then we find an edge in G[X] which connects them and creates the \mathcal{M} -alternating a_2b_2 -path P_2 , which uses at most $2\log_2 n$ vertices of H. We then remove all the vertices from the trees which do not correspond to P_2 in a leaf-by-leaf manner. As such, Lemma 3.16 implies that the resulting forest is a good embedding. We continue this operation for every pair in \mathcal{P} ; note that we can indeed do this since every \mathcal{M} -alternating path we create uses at most $2\log_2 n$ vertices of H and so, at any given point of the process the current forest to which we are applying Theorem 3.15 is of size at most $|\mathcal{P}| \cdot 2\log_2 n + 2 \cdot |V(H)|/50 \le |\mathcal{M}|/50 + |V(H)|/25 < |V(H)|/10$, where the first term is an upper bound on the total number of vertices of H used in the paths P_1, P_2, \ldots and the second upper bounds the number of vertices in the current binary out-oriented trees which are being used - and thus, this theorem can indeed be applied. This completes the proof since the procedure constructs the desired vertex-disjoint \mathcal{M} -alternating paths in $G \cup \mathcal{M}$.

4 Robust Pósa rotation

In this section we will give various tools which allow us to close paths into cycles. More specifically, we will say that a path P can be closed into a cycle if there exists a cycle on the same vertex set V(P). Our approach to do this heavily relies on the so called rotation-extension technique, invented by Pósa in [44], which has subsequently been applied to several problems related to Hamiltonicity (see e.g., [13, 27, 33]). We will in turn develop new tools in order to use this technique in a robust manner. First, we need some definitions and initial lemmas.

4.1 Standard tools

We start with some definitions. Let $P = (v_1, v_2, ..., v_l)$ be a path in a graph G. For a vertex $u = v_i \in V(P)$, we define, as usual, $N_P(u) := \{v_{i-1}, v_{i+1}\}$ to be the set of neighbours of u in P. When necessary and when such an ordering of the path is defined, we will use $u_P^+ := v_{i+1}$ to denote its right neighbour in P and $u_P^- := v_{i-1}$ to denote its left neighbour. Given a subset $X \subseteq V(P)$, we define $\inf_P(X)$ (or just $\inf_{i=1}^{N} X_i \in X_i$) when P is clear from the context) to be the set of vertices $v_i \in X$ with $i \notin \{1, l\}$ such that $v_{i-1}, v_{i+1} \in X$.

Definition 4.1. For two paths P, P' on the same vertex set V, we define the difference between P and P', denoted as dif(P, P'), to be the set of vertices $u \in V$ such that $N_P(u) \neq N_{P'}(u)$.

It is worth making some remarks about the above definition. First, note that clearly the set $V \setminus \operatorname{dif}(P, P')$ spans the same collection of vertex-disjoint paths in both P and P' - moreover, there are at most $|\operatorname{dif}(P, P')| + 1$ of these paths. Note also that for any subset $X \subseteq V \setminus \operatorname{dif}(P, P')$, we have $\operatorname{int}_P(X) = \operatorname{int}_{P'}(X)$. We can now start introducing the influential rotation-extension technique of Pósa.

Definition 4.2. If 1 < i < l and $v_i v_l$ is an edge of G, the path

$$P' = (v_1, v_2, \dots, v_i, v_l, v_{l-1}, \dots, v_{i+1})$$

is said to be a rotation of P with fixed endpoint v_1 , pivot v_i and broken edge v_iv_{i+1} . More generally, given a subset $X \subseteq V(P)$, we will say that a path P' is an (X,t)-rotation of P with fixed endpoint v_1 if it is the result of at most t consecutive rotations starting with P, each with fixed endpoint v_1 , broken edges in E(P) and pivots belonging to the set $\inf_{P}(X)$. Clearly, such a path P' has as endpoints the vertex v_1 and a vertex in X. We say that it is just a t-rotation (or respectively, an X-rotation) if there is no restriction on where the pivots and broken edges belong to (or respectively, on the number of rotations).

Again, we make two remarks about the above definition.

Observation 4.3. Let P' be a 1-rotation of P as above. Then, only the vertices v_i, v_{i+1}, v_l change neighbourhoods, that is, $N_P(u) \neq N_{P'}(u) \Rightarrow u \in \{v_i, v_{i+1}, v_l\}$. Moreover,

- For all $X \subseteq V(P)$ such that $v_l \notin X$, then either the new endpoint v_{i+1} of P' belongs to X or $\operatorname{int}_{P'}(X) = \operatorname{int}_P(X)$.
- If P' is a t-rotation of P, then $|\operatorname{dif}(P, P')| \leq 3t$.

Now, we give two lemmas. The first is an analogue of Lemma 3.11 about finding clean subsets in the context of paths and Pósa rotation. Let us first define such a subset.

Definition 4.4. Let G be an (n, d, λ) -graph and P a path in G. A subset $S \subseteq V(P)$ is said to be (P, δ) -clean if for all $v \in S$ we have that $|N(v) \cap \operatorname{int}_P(S)| \geq \delta$. We say only that S is P-clean if it is (P, d|S|/4n)-clean.

Given this, we now have the following analogous statement to Lemma 3.11.

Lemma 4.5. Let G be an (n, d, λ) -graph, P a path and A a subpath of P. If $|A| \ge 10\lambda n/d$, there is a P-clean subset $A' \subseteq A$ with $|A'| \ge 0.9|A|$.

Proof. Similarly to the proof of Lemma 3.11, we do the following process. Start with $A_0 := A$ and at each step $i \ge 1$, remove a vertex $u_i \in A_{i-1}$ such that $|N(u_i) \cap \operatorname{int}_P(A_{i-1})| < d|A|/4n$, setting $A_i := A_{i-1} \setminus \{u_i\}$; if no such vertex exists then stop the process. Clearly, if the process stops at some step $i \le 0.1|A|$, then the desired subset $A' := A_i$ is constructed. Suppose then, for sake of contradiction, that the process only stops after that and consider the set A_i with i := 0.1|A| + 1. Define also the set $U := \{u_1, \ldots, u_{i-1}\}$ which has size 0.1|A| and note that by construction, we must have that every vertex in U has at most d|A|/4n neighbours in $\inf_P(A_{i-1})$. In turn, we claim that this contradicts the first part of Lemma 3.11 with $S := \inf_P(A_{i-1})$. Indeed, note that since A is a subpath of P and we have $|A_{i-1}| \ge 0.9|A|$, it must be that $|S| = |\inf_P(A_{i-1})| \ge 0.7|A|$. In particular, d|A|/4n < d|S|/2n and so, the first part of Lemma 3.11 gives that

$$|A| \cdot |U| \le |A| \cdot \left(4\lambda^2 n^2 / d^2 |S|\right) \le 8\lambda^2 n^2 / d^2$$

which is a contradiction since $|A| \cdot |U| \ge 0.1 |A|^2 \ge 10\lambda^2 n^2/d^2$ by assumption.

The next lemma is a standard version of the usual Pósa rotation lemma which has been implicitly used for several Hamiltonicity problems. Its proof can be found, e.g., in [33].

Lemma 4.6 ([33]). Let G be an (n, d, λ) -graph with $d \ge 10\lambda$, P a maximal path in G with endpoints x, y and $|P| \ge n/10$. Let $Z_1 \subseteq Z_2 \subseteq \ldots \subseteq V(P)$ be a nested sequence defined as follows: for each $i \ge 1$, Z_i is the set of vertices z for which there exists an i-rotation of P with fixed endpoint y which is a zy-path. Then, $Z_1 \ne \emptyset$ and $|Z_{i+1}| \ge \min(2|Z_i|, n/100)$ for all $i \ge 1$.

We remark that a proof of this will be implicit in the more general Lemma 4.8 in the next section. Crucially, the above lemma implies the following.

Corollary 4.7. Let G be an (n, d, λ) -graph with $d \ge 10\lambda$, P a maximal path in G with endpoints x, y and $|P| \ge n/10$. Let $R \subseteq V(P)$ be a set such that $|\operatorname{int}_P(R)| \ge 100\lambda^2 n/d^2$. Then, there exists a $(\log n)$ -rotation P' of P with y as one endpoint and the other endpoint $x' \in R$ such that $|\operatorname{int}_P(R) \setminus \operatorname{int}_{P'}(R)| \le 3$.

Proof. Let us apply Lemma 4.6 and take I to be the minimal i such that $|Z_i| \geq n/100$ - by the lemma, we have that $I \leq \log n - 1$. If $R \cap Z_I \neq \emptyset$, then let $x' \in R \cap Z_I$ belong to Z_j with j minimal. Then, we claim that x' satisfies the desired property. Indeed, it is such that there is a j-rotation P' of P with y as one endpoint and the other endpoint x'. Moreover, since j is minimal note that the endpoints created by each one of the previous rotations forming P' are not in R except for x'. Thus, by the first item of Observation 4.3, the only change to $\inf_{P}(R)$ occurs in the last step and so, by the same observation, it only affects at most three vertices, hence $|\inf_{P}(R) \setminus \inf_{P'}(R)| \leq 3$, as desired.

Suppose now that $R \cap Z_I = \emptyset$. Since $|Z_I||\text{int}_P(R)| \ge \lambda^2 n^2/d^2$, part (4) of Lemma 3.8 implies that there is an edge zv with $z \in Z_I$ and $v \in \text{int}_P(R)$. Now, consider the zy-path P_z which is an I-rotation of P with fixed endpoint y. Since $R \cap Z_I = \emptyset$, Observation 4.3 implies that $\text{int}_P(R) = \text{int}_{P_z}(R)$. Therefore, $v \in \text{int}_{P_z}(R)$ and so, the edge zv ensures that there is a path P' which is a rotation of P_z with fixed endpoint y and an x'y-path for some $x' \in R$. We now check that P', x' satisfy the desired conditions. Indeed, P' is an (I+1)-rotation of P with $I+1 \le \log n$. Further, since $\text{int}_P(R) = \text{int}_{P_z}(R)$ and P' is a 1-rotation of P_z we have that $|\text{int}_P(R) \setminus \text{int}_{P'}(R)| \le 3$.

4.2 The main technical rotation lemma

We will further need a much more robust extension of the Pósa rotation technique than Lemma 4.6. Informally, given a path P we will first want to perform rotations only in a certain *clean* subset of P and secondly, we want to be able to do this robustly, that is, possibly after some changes to the path have been done. Clearly, the challenge here is that these changes might alter the *clean* subset we wish to rotate in.

Lemma 4.8. Let G be an (n, d, λ) -graph with $d \ge 100\lambda$, P a path in G, $X, Y \subseteq V(P)$ be non-empty with X being a (P, δ) -clean subset for some $20\lambda \le \delta \le d$. Let also P' be another path with V(P') = V(P) and let x, y be its endpoints with $x \in X$. Suppose the following holds in the case that $P \ne P'$.

• At most $\delta/2$ vertices $w \in \operatorname{int}_P(X) \setminus \operatorname{int}_{P'}(X)$ are such that $\operatorname{dist}_G(w,x) \leq 2 \log |\operatorname{dif}(P,P')|$.

Then there are at least $\delta n/200d$ vertices $z \in X$ for which there exists a zy-path P_z which is an $(X, \log n)$ rotation of P' with fixed endpoint y and such that $|\operatorname{dif}(P_z, P') \cap Y| \leq 3 \log |Y|$.

Proof. Define $F := \operatorname{dif}(P, P')$, so that we have for all $u \notin F$ that $N_P(u) = N_{P'}(u)$, and as observed before, $\operatorname{int}_P(X \setminus F) = \operatorname{int}_{P'}(X \setminus F)$. Moreover, clearly we have $|\operatorname{int}_P(X \setminus F)| \ge |\operatorname{int}_P(X)| - 3|F|$ which implies the following: for every $S \subseteq X$ we have $|N(S) \cap \operatorname{int}_{P'}(X)| \ge |N(S) \cap \operatorname{int}_{P'}(X \setminus F)| \ge |N(S) \cap \operatorname{int}_P(X)| - 3|F|$. Now, for each $i \ge 0$, let us define $Z_i \subseteq X$ to be the set of vertices z for which there exists an (X, i)-rotation P_z of P' with fixed endpoint y which is a zy-path and such that $|\operatorname{dif}(P_z, P') \cap Y| \le 3\log |Y|$. Clearly, note that $\{Z_i\}_i$ is a family of nested sets with $Z_i \subseteq Z_{i+1}$ for each i. Further, the following recurrence holds.

Claim. For each $i < \log |Y|$, we have $|Z_{i+1}| \ge \frac{1}{2} |N(Z_i) \cap \text{int}_{P'}(X)| - \frac{3}{2} |Z_i|$. For $i \ge \log |Y|$, we have

$$|Z_{i+1}| \ge \frac{1}{2} |N(Z_i) \cap \operatorname{int}_{P'}(X)| - \frac{3}{2} |Z_i| - \frac{3}{2} |Y|.$$

Proof. Let us suppose first that $i < \log |Y|$. Fix an initial ordering $(v_1 := x, v_2, \dots, v_{l-1}, v_l := y)$ of the path P' and define the set of indices $J := \{1 < j < l : v_j \in N(Z_i) \cap \operatorname{int}_{P'}(X), v_{j-1}, v_j, v_{j+1} \notin Z_i\}$. Clearly

we have that $|J| \geq |N(Z_i) \cap \operatorname{int}_{P'}(X)| - 3|Z_i|$. Consider a vertex v_j with $j \in J$. We claim that one of v_{j-1}, v_{j+1} belongs to Z_{i+1} , thus showing that $|Z_{i+1}| \geq |J|/2$, which by the previous lower bound on |J| implies the desired recurrence. Indeed, note that since $v_j \in N(Z_i) \cap \operatorname{int}_{P'}(X)$, it has a neighbour $z \in Z_i$. By definition, there exists an (X,i)-rotation of P' with fixed endpoint y and whose other endpoint is z - let us denote this path by Q. Note that the subpath $v_{j-1}v_jv_{j+1}$ is still present in Q (albeit in maybe a different direction). Indeed, during the repeated rotation process of obtaining Q, by definition every broken edge must be incident to some vertex in Z_i and we have $v_{j-1}, v_j, v_{j+1} \notin Z_i$. Therefore, we can use the edge zv_j to perform a rotation of Q with fixed endpoint y, pivot v_j and broken edge either $v_{j-1}v_j$ or v_jv_{j+1} . This results in an (X,1)-rotation of Q and so, an (X,i+1)-rotation of P' with fixed endpoint y. The other endpoint of the path will be either v_{j-1}, v_{j+1} and so one of these two vertices belongs to Z_{i+1} . Note further that this path, which we denote as Q', an (X,i+1)-rotation of P' with fixed endpoint y, by Observation 4.3 is such that $|\operatorname{dif}(Q',P') \cap Y| \leq |\operatorname{dif}(Q',P')| \leq 3i+3 \leq 3 \log |Y|$.

For the case that $i \geq \log |Y|$, the proof is analogous, the only difference being that we define $J := \{1 < j < l : v_j \in N(Z_i) \cap \operatorname{int}_{P'}(X), v_{j-1}, v_j, v_{j+1} \notin Z_i \cup Y\}$, which has size at least $|N(Z_i) \cap \operatorname{int}_{P'}(X)| - 3|Z_i| - 3|Y|$. We can then show similarly as above that for each $v_j \in J$, one of v_{j-1}, v_{j+1} belongs to Z_{i+1} - and thus, $|Z_{i+1}| \geq |J|/2$. Indeed, there is a $z \in Z_i$ which is a neighbour of v_j , so that there is an (X,i)-rotation Q of P' with fixed endpoint y and whose other endpoint is z. The edge zv_j can then be used to extend this to an (X,i+1)-rotation of P' with fixed endpoint y and one of v_{j-1},v_{j+1} as the other endpoint - without loss of generality, assume that it is v_{j-1} and let this path be denoted as Q'. Note that $\operatorname{dif}(Q',P')\setminus\operatorname{dif}(Q,P')\subseteq\{v_j,v_{j-1}\}$ and so, since $v_{j-1},v_j,v_{j+1}\notin Y$, we have by definition that $|\operatorname{dif}(Q',P')\cap Y|\leq |\operatorname{dif}(Q,P')\cap Y|\leq 3\log|Y|$ as desired.

Now, note that by assumption we have $|Z_1| = |N(x) \cap \operatorname{int}_{P'}(X)| \ge \delta/2$. Indeed, since X is (P, δ) -clean, we have that $|N(x) \cap \operatorname{int}_P(X)| \ge \delta$ - therefore, if P = P' this trivially holds and if $P \ne P'$, then the assumption of the lemma implies that then $|N(x) \cap \operatorname{int}_{P'}(X)| \ge \delta/2$ (since any neighbour $w \in N(x)$ clearly has $\operatorname{dist}(w,x) \le 2\log|F|$). Thus, in order to finish the proof of the lemma we only need to show that while $|Z_i| \le \delta n/200d$, we have that $|N(Z_i) \cap \operatorname{int}_{P'}(X)| \ge 10|Z_i|$. Indeed, if that is the case, then by the claim above we have first for each such $\log |Y| > i \ge 1$ that $|Z_{i+1}| \ge 5|Z_i| - \frac{3}{2}|Z_i| \ge 2|Z_i|$. In particular, this implies that $|Z_I| \ge 2^I \ge |Y|$ for $I := \log |Y|$. Then, again by the claim above we have that for all $i \ge I$, $|Z_{i+1}| \ge 5|Z_i| - \frac{3}{2}|Z_i| - \frac{3}{2}|Y| \ge 5|Z_i| - 3|Z_i| \ge 2|Z_i|$. Therefore, $|Z_l| \ge \delta n/200d$ will occur for some $l \le \log n$. This in turn gives us the desired vertices z - those contained in Z_l .

Now, suppose then that i is such that $|Z_i| \leq \delta n/200d$ and consider $N(Z_i) \cap \operatorname{int}_{P'}(X)$. Note first that when $P \neq P'$ and $i < \log |F|$, since Pósa rotation implies that $N(Z_i)$ is composed of vertices which are at a distance at most $2i+1 < 2\log |F|$ of x, we have by assumption of the lemma that every vertex in Z_i has at least $\delta/2$ neighbours in $\operatorname{int}_{P'}(X)$ (recall that every such vertex has at least δ neighbours in $\operatorname{int}_{P}(X)$). Applying now Lemma 3.9 with G' := G, $A := Z_i$, $B := \operatorname{int}_{P'}(X)$ we get that $|N(Z_i) \cap \operatorname{int}_{P'}(X)| \geq 10|Z_i|$. Now, let us consider when $i \geq \log |F|$ or P = P' (then clearly $F = \emptyset$). Note that the previous observations imply that $|Z_i| \geq 2^{\log |F|} \geq |F|$ - hence, $|N(Z_i) \cap \operatorname{int}_{P'}(X)| \geq |N(Z_i) \cap \operatorname{int}_{P}(X)| - 3|Z_i|$ (note that this also holds if P = P'). Since every vertex in Z_i has at least δ neighbours in $\operatorname{int}_{P}(X)$, Lemma 3.9 then implies also that $|N(Z_i) \cap \operatorname{int}_{P}(X)| \geq \min\left(\frac{\delta^2}{8\lambda^2}, 20\right)|Z_i| \geq 13|Z_i|$. Hence, $|N(Z_i) \cap \operatorname{int}_{P'}(X)| \geq 13|Z_i| - 3|Z_i| \geq 10|Z_i|$, as desired.

Note that as a corollary, we get the following which tells us that we can always rotate inside P-clean subsets and avoid certain vertices.

Corollary 4.9. Let G be an (n,d,λ) -graph with $d \geq 100\lambda$, P a path in G with endpoints x,y and $X,Y \subseteq V(P)$ be such that X is a (P,δ) -clean subset for some $20\lambda \leq \delta \leq d$ and $x \in X$. Then there are at least $\delta n/200d$ vertices $z \in X$ for which there exists a zy-path P_z which is an $(X, \log n)$ -rotation of P with fixed endpoint y and such that $|\operatorname{dif}(P_z, P) \cap Y| \leq 3\log |Y|$.

4.3 A method for closing paths into cycles

We finish with a general setup which will always allow one to close a given path into a cycle. This was already alluded to in Section 2. The challenges in achieving this were also discussed there and will be dealt with in the next section. Before stating the setting, let us give a definition. For a path P and a subset $A \subseteq V(P)$, the graph induced by P on A, that is, P[A], is a disjoint union of paths - we will let $c_P(A) := c(P[A])$ denote the number of such paths, that is, the number of components of P[A].

Lemma 4.10. Let G be an (n,d,λ) -graph, P a path in G with endpoints x,y and $A,B \subseteq V(P)$ be disjoint sets such that the following holds. Both A and B contain $(P,200\lambda \cdot 2^{\min(c_P(A),c_P(B))})$ -clean subsets $A' \subseteq A, B' \subseteq B$ with $x \in A'$ and $y \in B'$. Then, P can be closed into a cycle.

Proof. Let us fix an ordering $P = (x, v_2, \dots, v_{l-1}, y)$ and suppose without loss of generality that $c_P(A) \ge c_P(B)$. First, we apply Corollary 4.9 with $X := A', Y := \emptyset$ thus finding at least $(\lambda n/d) \cdot 2^{c_P(B)}$ vertices z such that there is a zy-path P_z which is an A'-rotation of P with fixed endpoint y - let Z_1 denote the set of these vertices. Note crucially that for each $z \in Z_1$, since P_z is an A'-rotation of P, we have that $P_z[B] = P[B]$ and so, $c_{P_z}(B) = c_P(B)$. Moreover, the following holds.

Claim. There exists a subset $Z_1' \subseteq Z_1$ of size at least $|Z_1| \cdot 2^{-c_P(B)}$ such that the following holds: for all $v \in B$ and $z, z' \in Z_1'$ we have $v_{P_z}^+ = v_{P_{z'}}^+$ and $v_{P_z}^- = v_{P_{z'}}^-$.

Proof. For the purpose of the argument let us fix $c := c_P(B)$. As noted above, we have for each $z \in Z_1$ that $P_z[B] = P[B]$. Indeed, the only difference between B in P_z and in P is the direction in which each path of $P_z[B] = P[B]$ is traversed. More precisely, let us recall the ordering of P as (x, \ldots, y) , that is, from x to y and also consider the ordering of P_z from z to y. Let Q be one of the c sub-paths of P which form B, let u, v be its endpoints and assume that in the given ordering of P we have that Q is ordered from u to v. We know that Q is also a sub-path of P_z , but we are unaware of which direction Q is ordered in the ordering of P_z from z to y, there being 2 possibilities $(u \to v \text{ or } v \to u)$. Furthermore, this direction will determine $v_{P_z}^+$ and $v_{P_z}^-$ for all $v \in Q$. Therefore, by pigeonholing, since there are c sub-paths Q we can find a fraction of 2^{-c} of the vertices $z \in Z_1$ for which all of these are equal, as desired.

Note that the above claim implies that if a vertex $w \in B$ is such that for some $z \in Z_1'$ there is a zw-path which is a B'-rotation of P_z with fixed endpoint w, then this is the case for all other vertices $z' \in Z_1'$. Therefore, letting $W \subseteq B$ denote the set of such vertices w we have that for all $z \in Z_1'$, $w \in W$, there is a zw-path on the vertex set V(P). Hence, if there exists an edge zw between two such vertices, then P can be closed into a cycle. In order to find such an edge, we only need to show that $|W| \ge \lambda n/d$ since then $|W||Z_1'| \ge (\lambda n/d)|Z_1| \cdot 2^{-c_P(B)} \ge (\lambda n/d)^2$ and thus, part (4) of Lemma 3.8 guarantees that the edge exists. To show that $|W| \ge \lambda n/d$, fix some $z \in Z_1'$ and note that B' is a $(P_z, 200\lambda)$ -clean subset of P_z and $y \in B'$. So, Corollary 4.9 applied on the path P_z with X := B' implies that $|W| \ge \lambda n/d$.

5 Hamilton cycles with robust Pósa rotation

In this section we will prove Theorem 1.2. In order to do this, we will show that it will always be possible to close a maximal path into a cycle. Clearly, if a graph G is connected (as are (n, d, λ) -graphs with $\lambda < d/2$ by part (5) of Lemma 3.8) and every maximal path can be closed into a cycle, then it is Hamiltonian. As expected, our arguments here will build on the tools given in the previous section. First, we introduce the notion of a *clean* collection of intervals of a path. This will be crucial for then closing this path into a cycle.

Definition 5.1. Given a path P, a collection $Q = \{Q_1, \ldots, Q_k\}$ of disjoint sub-paths of P is said to be (δ, γ, k) -clean for P if it is equipped with a subset $S \subseteq \bigcup_i Q_i$ and has the following properties.

- 1. $|Q_i| \ge 0.99 |P|/k$ and $|S \cap Q_i| \ge 0.99 |Q_i|$ for each *i*.
- 2. Every vertex $v \in S$ has at least δ neighbours in $\operatorname{int}_P(S)$ that is, S is (P, δ) -clean.
- 3. For each vertex $v \in S$ there are at least γk indices j such that $|N(v) \cap \operatorname{int}(S \cap Q_i)| \geq \delta/k$.

We remark first that the reader should think of the above definition with δ being less than d|S|/20n, rather than d|S|/4n, where the second condition would tell us that S is P-clean. This then makes the third condition much more plausible to obtain. We remark also that this notion appeared in [30] in the context of random graphs in a much more simplified form. We can now show that every sufficiently large path in a pseudorandom graph contains a clean collection - we remark that in applications, the value of k in the lemma below will be small, of polylogarithmic order in terms of n.

Lemma 5.2. Let $\alpha := 10^{-8}$ and G be an (n, d, λ) -graph with $\lambda \leq \alpha d$ and P a path in G of size at least n/3. Then for all k there is a $\left(d/100, \frac{\alpha d^2}{k\lambda^2}, k\right)$ -clean collection for P.

Proof. Let $\delta = d/100$, $\beta := 10^{-5}$ and $\gamma := \frac{\alpha d^2}{k\lambda^2}$. Let us first partition P into disjoint subpaths I_1, \ldots, I_r each of size |P|/r with r := 1.01k. Initially take S := V(P) and iteratively do the following procedure until it is no longer possible: (1) If there exists an index j such that $|S \cap I_j| < 0.99|I_j|$, then remove every vertex of I_j from S and continue; (2) If that is not the case and there is a vertex $v \in S$ with degree less than δ in int(S), remove it from S and continue; (3) If that is also not the case and there is a vertex $v \in S$ such that there are less than γr indices j for which $|N(v) \cap int(S \cap I_j)| \ge \delta/k$, remove it from S and continue.

The goal is then to show that this process must stop while the number of indices j such that $|S \cap I_j| \ge 0.99|I_j|$ is at least k - we then take $\mathcal Q$ to be the collection of these sub-paths and it will have the desired properties. Suppose indeed that this is not the case. Then, we can assume that we achieve a stage of the process when the number of indices j such that $|S \cap I_j| \ge 0.99|I_j|$ is precisely k - let $\mathcal I$ be the set of such indices. Define also $R \subseteq V(P) \setminus S$ to be the set of vertices which have been previously removed in the process because one of the options (2) or (3) described before occurred. Observe first then that the following must hold.

Claim. $|R| \ge (r - k) \cdot 0.01 |P|/r \ge \beta n$.

Proof. By the definition of the process, it must be that every subpath I_j that has been removed from S (that is, using step (1)) is such that $0.01|I_j| = 0.01|P|/r$ of its vertices were previously removed individually and thus, are contained in R. Since by assumption exactly r - k subpaths I_j have been removed from S, we have that $|R| \ge (r - k) \cdot 0.01|P|/r \ge \beta n$.

Moreover, by definition every vertex $v \in R$ satisfies one of the following properties.

- v has degree less than δ in int(S).
- There are fewer than γr indices $j \in \mathcal{I}$ for which $|N(v) \cap \operatorname{int}(S \cap I_j)| \geq \delta/k$.

Let $R_1, R_2 \subseteq R$ respectively denote the sets of vertices of R satisfying these properties. We will now give upper bounds for both R_1, R_2 using Lemma 3.8 and thus achieve a contradiction to the above claim. First, consider R_1 . Note that since $|S| \ge \sum_{j \in \mathcal{I}} |S \cap I_j| \ge \sum_{j \in \mathcal{I}} 0.99 |I_j| \ge k \cdot 0.99 |P|/r \ge 0.98 |P|$, we have that $|\operatorname{int}(S)| \ge |S| - 2|P \setminus S| \ge 0.9|P| \ge n/4$. Therefore, $\delta \le d|\operatorname{int}(S)|/2n$ and so, by Lemma 3.11, we must have $|R_1| \le 4\lambda^2 n^2/d^2 |\operatorname{int}(S)| \le \beta n/3$.

Now, consider R_2 . We claim that it has size at most $\beta n/3$, which is a contradiction since then $\beta n \leq |R| \leq |R_1| + |R_2| \leq 2\beta n/3$. Indeed, define an auxiliary bipartite graph H with one part A corresponding to R_2 and the other part B corresponding to \mathcal{I} - put an edge in H between $v \in A$ and

 $j \in B$ if $|N(v) \cap \operatorname{int}(S \cap I_j)| < \delta/k$. By definition of R_2 , every vertex in A has more than $k - \gamma r \ge (1 - 2\gamma)k = (1 - 2\gamma)|B|$ neighbours in B and so, we can greedily find a subset $B' \subseteq B$ of $\frac{1}{4\gamma}$ vertices with at least |A|/2 common neighbours in A. Let A' denote this set of common neighbours and let $X := \bigcup_{j \in B'} \operatorname{int}(S \cap I_j)$. Recall that for each $j \in B'$ we have $|S \cap I_j| \ge 0.99|I_j|$ and so, $\operatorname{int}(S \cap I_j) \ge |S \cap I_j| - 2|I_j \setminus S| \ge 0.97|I_j| \ge 0.9|P|/r \ge n/5r$ - hence, $|X| \ge |B'|n/5r \ge n/20\gamma r$. Further, by definition of H we must have that every vertex $v \in A'$ has at most $\delta |B'|/k = \delta/4\gamma k < d|X|/2n$ neighbours in X. So, by Lemma 3.11 we must have that $|R_2| \le 2|A'| \le 8\lambda^2 n^2/d^2|X| < \beta n/3$.

5.1 Proof of Theorem 1.2

We will now prove the first main theorem of the paper. Let G be an (n, d, λ) -graph with $d/\lambda \geq C(\log n)^{1/3}$ where C is some large absolute constant. Fix $\delta = d/100$, $\gamma := \frac{C}{(\log n)^{1/3}} \geq C^2 \lambda/d$, $k := 30 \log n$ and $D = d/(\log n)^{1/3}$. Let P be a path of maximal length in G - we will show that P can be closed into a cycle. Since G is connected this clearly implies that G is Hamiltonian. Note also that an undirected version of Lemma 3.13 along with part (4) of Lemma 3.8 gives that $|P| \geq n - 2\lambda n/d = n - o(n)$.

Let us also denote P by an ordering $v_1v_2 \dots v_{l-1}v_l$, with $x := v_1, y := v_l$ as its endpoints. The first step is to apply Lemma 5.2. Indeed, we partition P into two subpaths P_1, P_2 of size |P|/2. By Lemma 5.2, both paths have (δ, γ, k) -clean collections - indeed note that $\frac{10^{-8}d^2}{k\lambda^2} \ge \frac{C^2}{30 \cdot 10^8 (\log n)^{1/3}} \ge \gamma$ if C is large enough. Let $Q^1 = \{Q_1^1, \dots, Q_k^1\}, Q^2 = \{Q_1^2, \dots, Q_k^2\}$ be the given collection of disjoint sub-paths of P_1, P_2 respectively, and $S_1 \subseteq \bigcup_i Q_i^1, S_2 \subseteq \bigcup_i Q_i^2$ be the subsets with the given properties - note that $|\inf_P(S_1)|, |\inf_P(S_2)| \ge 0.4|P|$ and that S_1, S_2 are disjoint.

First, we note that we can use Lemma 4.6 in order to pass to a path P' which has special properties, one of them being that one of its endpoints belongs to S_1 and the other to S_2 .

Claim 5.3. There exists a path P' on the same vertex set as P with endpoints x', y' such that the following hold for some $a \in \{1, 2\}$ (let $b \in \{1, 2\}$ with $b \neq a$).

- $x' \in S_a, y' \in S_b$.
- $|\operatorname{dif}(P, P')| \leq 10 \log n \ and \ |\operatorname{int}_P(S_b) \setminus \operatorname{int}_{P'}(S_b)| \leq 9.$
- $\operatorname{int}_P(S_a) \setminus \operatorname{int}_{P'}(S_a)$ contains at most D/2 vertices z such that $\operatorname{dist}(z, x') \leq 2 \log \log n$.

Proof. Applying Corollary 4.7 directly to P and the set $R = S_1 \cup S_2$, we can find a $(\log n)$ -rotation P^1 of P with fixed endpoint y such that the other endpoint, say x', belongs to some S_a with $a \in \{1, 2\}$. Furthermore, we have that $|\inf_P(S_1 \cup S_2) \setminus \inf_{P^1}(S_1 \cup S_2)| \leq 3$. Similarly, we can again apply that corollary to P^1 and the set $R = S_1 \cup S_2$ to find a $(\log n)$ -rotation P^2 of P^1 with fixed endpoint x' such that the other endpoint, say z, belongs to $S_1 \cup S_2$. Furthermore, we have that $|\inf_P(S_1 \cup S_2) \setminus \inf_{P^2}(S_1 \cup S_2)| \leq 6$. Now, if $z \in S_b$, where $b \in \{1, 2\}$ and $b \neq a$, then we are done by letting y' := z and $P' := P^2$. If $z \in S_a$, we need still to apply Corollary 4.9 to P^2 .

Let us firstly consider the case that $\delta \geq 100(\log n)^{4/3}$. Note that since $|\operatorname{int}_P(S_1 \cup S_2) \setminus \operatorname{int}_{P^2}(S_1 \cup S_2)| \leq 6$, we have that S_a is $(P^2, \delta - 6)$ -clean and so, we can rotate the endpoint z and use Corollary 4.9 to conclude that there are at least n/10000 vertices $w \in S_a$ for which there exists a wx'-path P_w which is an $(S_a, \log n)$ -rotation of P^2 with fixed endpoint x'. Since $|\operatorname{int}_{P^2}(S_b)| \geq |P|/3 \geq n/4$, then property (4) of Lemma 3.8 implies that there is an edge between such a w and $\operatorname{int}_{P_2}(S_b)$ which in turn, implies that there is some $y' \in S_b$ and a y'x'-path P' which is a $(\log n+1)$ -rotation of P^2 with x' as a fixed endpoint. Note that all the conditions are satisfied: $y' \in S_b$; because of the assumption on δ we have that $|\operatorname{dif}(P, P')| \leq 10 \log n \leq D/2$ and so, the first part of the second condition and the third condition are satisfied; and since the path P_w was an S_a -rotation of P^2 and S_a , S_b are disjoint, we have that any change to the interior of S_b only occurred with the last rotation step resulting in P', hence $|\operatorname{int}_P(S_b) \setminus \operatorname{int}_{P'}(S_b)| \leq |\operatorname{int}_P(S_b) \setminus \operatorname{int}_{P^2}(S_b)| + 3 \leq 9$.

Now, let us consider when $\delta < 100(\log n)^{4/3}$ and so, $d < 10000(\log n)^{4/3}$ - recall that we always have $d \ge C^2(\log n)^{2/3}/2$ (since necessarily $d/(\log n)^{1/3} \ge C\lambda \ge C\sqrt{d/2}$ by the remark after Theorem 3.7) and thus, $D \ge (\log n)^{1/3}$. Since we want to ensure the last property of the statement, consider the set $Y \subseteq S_a$ consisting of all the vertices in S_a which are at a distance at most $2\log\log n$ from x' - by the assumption on d, we have that $|Y| \le d^{2\log\log n} < 2^{10(\log\log n)^2}$. We now apply Corollary 4.9 again, this time with the above set Y. Indeed, as previously, recall that S_a is $(P^2, \delta - 6)$ -clean and so, the corollary implies that there are at least n/10000 vertices $w \in S_a$ for which there exists a wx'-path P_w which is an $(S_a, \log n)$ -rotation of P^2 with fixed endpoint x' and such that

$$|\operatorname{dif}(P_w, P^2) \cap Y| \le 3\log|Y| \le 30(\log\log n)^2 \le D/2 - 6.$$

Since $|\inf_{P_2}(S_b)| \geq n/3$, then property (4) of Lemma 3.8 implies that there is an edge between such a w and $\inf_{P_2}(S_b)$. To conclude, this implies that there is some $y' \in S_b$ and a y'x'-path P' which is a $(\log n+1)$ -rotation of P^2 with x' as a fixed endpoint and such that $\operatorname{dif}(P', P^2) \cap Y$ has size at most D/2-6. Since we have that $|\inf_{P}(S_a) \setminus \inf_{P^2}(S_a)| \leq 6$ then it must be that $|(\inf_{P}(S_a) \setminus \inf_{P'}(S_a)) \cap Y| \leq D/2$, as desired. The other conditions can be checked as in the first case.

Let us define $F := \operatorname{dif}(P,P')$ so that by the second property of the claim above, we have $|F| \leq k/3$. Now, consider the given sub-paths Q_j^a, Q_j^b of P and note that these might not be sub-paths of P'. Clearly however, $P'[Q_j^a], P'[Q_j^b]$ are disjoint unions of paths - we will denote these components of $P'[Q_j^a]$ by $C(Q_j^a, 1), \ldots, C(Q_j^a, i_{a,j})$ and the components of $P'[Q_j^b]$ by $C(Q_j^b, 1), \ldots, C(Q_j^b, i_{b,j})$. We will refer to them as the parts of Q_j^a, Q_j^b (in P'), respectively. We will further say that Q_j^a is a broken sub-path if $i_{a,j} > 1$ and unbroken otherwise (and analogously for b). Note that we must have $2k \leq \sum_{1 \leq j \leq k} i_{a,j} + \sum_{1 \leq j \leq k} i_{b,j} \leq 2k + |F| + 1 \leq 7k/3 + 1$ and in particular, there are at most |F| + 1 parts of broken sub-paths.

The goal will be to achieve the setting of Lemma 4.10. For this, we will first prove the following claim with a standard averaging argument.

Claim 5.4. There exist two disjoint subpaths $A, B \subseteq P'$ such that both contain precisely $\gamma k/4$ parts of sub-paths in $Q_1 \cup Q_2$ and at least $\gamma k/8$ unbroken paths.

Proof. Let us first take all the paths $C(Q_j^a, l)$, $C(Q_j^b, l)$ and denote them by C_1, \ldots, C_r where the order is given by the order in which they appear in the path P' (for this, fix a direction of P' as $x' \to P' \to y'$). Recall that $r = \sum_{1 \le j \le k} i_{a,j} + \sum_{1 \le j \le k} i_{b,j} \in [2k, 7k/3 + 1]$. Clearly, we can use this order to define a cover of $\bigcup_t C_t$ with disjoint sub-paths of P' denoted by P'_1, P'_2, \ldots , such that each P'_i contains r_i consecutive sub-paths C_i for some $r_i \in [\gamma k/4, \gamma k/2]$ - by construction, there are at least $\frac{r}{r} \ge \frac{2r}{ck}$ paths P'_i .

sub-paths C_j for some $r_i \in [\gamma k/4, \gamma k/2]$ - by construction, there are at least $\frac{r}{\max_i r_i} \geq \frac{2r}{\gamma k}$ paths P'_i . Now, let us count how many paths P'_j do not satisfy the property in the statement. Note first that every sub-path P'_i has at least $\gamma k/4$ parts of sub-paths in $Q_1 \cup Q_2$ by construction. Then, if P'_j does not contain a sub-path $P''_j \subseteq P'_j$ satisfying the property, it must contain at least $\gamma k/8$ parts from broken sub-paths. In turn, recall that we observed that the number of such parts must be at most $|F| + 1 \leq k/3 + 1$ and so, since the sub-paths P'_i cover all parts of sub-paths in $Q_1 \cup Q_2$, we have that at most $\frac{|F|+1}{\gamma k/8} \leq 3/\gamma$ sub-paths P'_i do not satisfy the property. Concluding, since $3/\gamma < \frac{4}{\gamma} - 2 \leq \frac{2r}{\gamma k} - 2$ and there are at least $\frac{2r}{\gamma k}$ paths P'_i , we have the existence of the two desired subpaths A, B.

We are now almost in the setting where we can apply Lemmas 4.8 and 4.10 in order to finish the proof. Indeed, note first the following.

Claim 5.5. $S_a \setminus (A \cup B), S_b \setminus (A \cup B) \subseteq P$ are (P, D)-clean subsets.

Proof. We only prove this for S_a , since it is analogous for S_b . Let then $v \in S_a$ and recall that the third condition of Definition 5.1 implies that there exist at least γk many sub-paths Q_j^a such that |N(v)|

 $\operatorname{int}_P(S_a \cap Q_j^a)| \geq \delta/k$. Recall also that Claim 5.4 implies that each A and B contain precisely $\gamma k/4$ parts from sub-paths in \mathcal{Q}_a and so, there are at least $\gamma k/2$ sub-paths Q_j^a disjoint to $A \cup B$ so that $|N(v) \cap \operatorname{int}_P(S_a \cap Q_j^a)| \geq \delta/k$. Clearly, this implies that

$$|N(v) \cap \operatorname{int}_P(S_a \setminus (A \cup B))| \ge (\gamma k/2) \cdot (\delta/k) \ge D,$$

thus showing that $S_a \setminus (A \cup B)$ is (P, D)-clean.

Note now that by Claim 5.4, the given sub-paths A, B of P' contain both at least $\gamma k/8$ unbroken paths Q_i^a , each of which have size at least 0.99|P|/k and so, we must have that $|A|, |B| \ge (\gamma k/8) \cdot (0.99|P|/k) > \gamma |P|/10 \ge 10\lambda n/d$ (since $|P| \ge n - o(n)$). Therefore, Lemma 4.5 implies that there exist P'-clean subsets $A' \subseteq A, B' \subseteq B$ of size at least 0.9|A| and 0.9|B| respectively. Finally, this allows us to get to the following setting in which we can apply Lemma 4.10.

Claim 5.6. There is a path P'' on the same vertex set as P' with $|\operatorname{dif}(P', P'')| \leq \log n$ and the following properties.

- (1) The endpoints of P'' belong to A' and B'.
- (2) Both $A', B' \subseteq P''$ are $(P'', 800\lambda)$ -clean.
- (3) $c_{P''}(A), c_{P''}(B) \leq 2.$

Proof. For this, we apply Lemma 4.8. Indeed, let $X_a := S_a \setminus (A \cup B)$, $X_b := S_b \setminus (A \cup B)$ which are disjoint and by the previous claim are (P, D)-clean. From Claim 5.3 we also know that $|\inf_P(S_b) \setminus \inf_{P'}(S_b)| \leq 9$, which implies then that X_b is (P', D-9)-clean, and that $\inf_P(X_a) \setminus \inf_{P'}(X_a)$ contains at most D/2 vertices z such that $\operatorname{dist}(z, x') \leq 2 \log \log n$ - and note clearly (for a future application of Lemma 4.8) that $2 \log |F| \leq 2 \log \log n$.

Therefore, Lemma 4.8 implies firstly that there exist at least $\lambda n/d$ vertices $z \in X_a$ for which there exists a zy'-path which is an $(X_a, \log n)$ -rotation of P' with fixed endpoint y' (note that this lemma can be applied since $D = d/(\log n)^{1/3} \ge C\lambda$ and C is large enough). Let Z_a denote the set of such vertices z, so that $|Z_a| \ge \lambda n/d$. Since A is a sub-path of P', then $|\inf_{P'}(A')| \ge |A| - 3|A \setminus A'| \ge 0.7|A| \ge \lambda n/d$ and therefore, there exists an edge between Z_a and $\inf_{P'}(A')$. Crucially, by definition of Z_a this implies that there exists a vertex $x'' \in A'$ such that there is an $(X_a \cup A, \log n + 1)$ -rotation of P' with fixed endpoint y' which is an x''y'-path, which we shall call P'_a . Moreover, it must be that $c_{P'_a}(A) \le 2$, since the rotation only touches A at the last step. Finally, note also that $\inf_{P'_a}(P', P'_a) \cap (X_b \cup B) = \emptyset$ and in particular, B is still the same sub-path in P'_a .

Secondly, Corollary 4.9 applied with X_b implies that there exist at least $\lambda n/d$ vertices $z \in X_b$ for which there exists a x''z-path which is an $(X_b, \log n)$ -rotation of P_a' with fixed endpoint x''. Let Z_b denote the set of such vertices z, so that $|Z_b| \geq \lambda n/d$. Much like before, since $|\inf_{P_a'}(B')| = |\inf_{P'}(B')| \geq |B| - 3|B \setminus B'| \geq 0.7|B| \geq \lambda n/d$, there exists an edge between Z_b and $\inf_{P_a'}(B')$. Then, there is a vertex $y'' \in B'$ such that there is an $(X_b \cup B, \log n + 1)$ -rotation of P_a' with fixed endpoint x'' which is a x''y''-path - this will be our desired path P''. Note also that like before, we have $c_{P''}(B) \leq 2$ and that $\operatorname{dif}(P_a', P'') \cap (X_a \cup A) = \emptyset$, so that A is still the same in P'' as it was in P_a' , in particular, $c_{P''}(A) \leq 2$ as desired. Note that clearly properties (1),(3) are satisfied. For property (2) observe that since A', B' were P'-clean and $c_{P''}(A), c_{P''}(B) \leq 2$, it must be that every vertex $v \in A'$ (respectively, $v \in B'$) has at least $d|A'|/4n - 2 \geq \gamma d/100 \geq C^2\lambda/100 \geq 800\lambda$ (respectively, $d|A'|/4n - 2 \geq 800\lambda$) neighbours in $\inf_{P''}(A')$ (respectively, $\inf_{P''}(B')$), and so, they are $(P'', 800\lambda)$ -clean. Indeed, we used that $|A'| \geq 0.9|A|$, $|A| \geq \gamma |P|/10$, |P| = n - o(n) and $\gamma \geq C^2\lambda/d$.

We can now finish immediately by applying Lemma 4.10 to P''.

6 Linear forests with few paths and good endpoints

In this section we will consider the problem of finding (in pseudorandom graphs) spanning linear forests with two special properties - they have a small number of paths; and all these paths have *good* endpoints. As indicated in Section 2, this will be crucial in the proof of Theorem 1.3.

Proposition 6.1. Let G be an (n, d, λ) -graph, $X \subseteq Y \subseteq V(G)$ and $\delta \ge 10^4 \lambda$ be such that $|X| \le \delta n/1000d$ and both $Y, Y \setminus X$ are δ -clean. Then, there exists a spanning linear forest \mathcal{F} of G[Y] such that the following hold.

- The number of paths in \mathcal{F} is $O(n/\delta^{1/5})$.
- Every path in \mathcal{F} has its endpoints in $Y \setminus X$.

First, let us discuss the above statement in the case that $X = \emptyset$ - here we need only to find a spanning linear forest \mathcal{F} of G[Y] which does not have many components. This can be achieved as follows: Using the expansion properties of G, one can find a relatively dense regular subgraph of G[Y]. Invoking known results concerning the linear arboricity conjecture (see Theorem 6.2 below), there is a decomposition of this subgraph into an almost optimal number of linear forests, and simply by averaging, one of these must consist of few paths.

In general, if X is non-empty, we will first have to absorb the vertices of X into $Y \setminus X$, that is, we cover X with a linear forest whose paths have their endpoints outside X. The idea is then to contract the pairs of endpoints and find a linear forest with few paths in the new graph. By reversing the contractions, we want to insert the paths covering X into the paths of this linear forest. To make this feasible, we will have to work in a directed setting. Let us then first recall the following classical result of Alon [3] (see also [7]) on the linear arboricity of regular digraphs, which will be crucial for us to find a spanning linear forest with few paths. By a r-regular digraph we mean a digraph in which every vertex has in-degree and out-degree r.

Theorem 6.2. Every r-regular digraph can be covered by $r + O(r^{4/5})$ linear forests.

We will now apply this theorem to find a linear forest with few components in digraphs with suitable expansion properties. (We will later see that a random orientation of G[Y] satisfies these properties.)

Lemma 6.3. Let H be an n-vertex digraph and r be such that the following holds. For every subdigraph $H' \subseteq H$ with $\Delta^+(H'), \Delta^-(H') < r$, we have that for all subsets $S \subseteq V(H)$,

$$|\Gamma_{H\backslash H'}^{+}(S)| \ge |S|.$$

Then, H contains a spanning linear forest \mathcal{F} with $O\left(n/r^{1/5}\right)$ paths.

Proof. In order to construct \mathcal{F} , we will only need to show that H contains an r-regular spanning subdigraph. Indeed, by Theorem 6.2, such a digraph can be covered by $r + O(r^{4/5})$ linear forests. One of the linear forests in such a covering must contain at least $rn/(r + O(r^{4/5}))$ edges and hence consist of at most $O(n/r^{1/5})$ paths, and therefore, H will contain such a linear forest.

Note we can construct the r-regular spanning subdigraph by finding r edge-disjoint 1-regular subdigraphs and therefore it is sufficient to show that for any $H' \subseteq H$ with $\Delta^+(H'), \Delta^-(H') < r$, the digraph $H \setminus H'$ contains a 1-regular spanning subdigraph. For this, we define an auxiliary bipartite graph $B = (V_1, V_2)$ with both parts V_1, V_2 being copies of V(H); a pair xy with $x \in V_1, y \in V_2$ is an edge in B if $x \to y$ is an edge of $H \setminus H'$. Now note that B contains a perfect matching, since for every $S \subseteq V_1$, we have that $|N_B(S)| = |\Gamma^+_{H \setminus H'}(S)| \ge |S|$ by assumption, and therefore Theorem 3.5 applies. To conclude, note that a perfect matching in B corresponds precisely to a 1-regular spanning subdigraph of $H \setminus H'$.

The next lemma deals with the absorption of the vertices of X. Moreover, its proof contains a novel covering idea which we believe may have other applications.

Lemma 6.4. There exists a linear forest \mathcal{F}' in G[Y] which covers X and such that every path has its endpoints in $Y \setminus X$ and all its inner vertices in X.

Proof. The main idea is, instead of finding \mathcal{F}' directly, we observe that it is sufficient to show the existence of a forest \mathcal{F} in G[Y] consisting of disjoint binary trees which all together cover X and such that all the leaves are contained in $Y \setminus X$ and all the non-leaves in X. Indeed, we can from \mathcal{F} find the desired linear forest \mathcal{F}' by doing the following operation. Consider some binary tree T in \mathcal{F} and take a path P in T which contains the root of T and whose endpoints are leaves (observe that P exists since the root of T has degree two in T). Note that $\mathcal{F} - V(P)$ is still a disjoint union of binary trees and it covers $X \setminus V(P)$. Thus, we can repeat the same operation until all of X has been covered by these paths, which altogether form the desired linear forest \mathcal{F}' (we can in every step discard the vertices of the forest which become isolated, since those must belong to $Y \setminus X$).

Now, in order to find the forest \mathcal{F} let us note that there exists an ordering x_1, x_2, \ldots, x_t of the vertices of X such that for all i, we have

$$|N_G(x_i) \cap (\{x_1, \dots, x_{i-1}\} \cup (Y \setminus X))| > \delta/2.$$

Indeed, suppose that we have such a partial sequence x_1, \ldots, x_{i-1} . If no candidate for the vertex x_i exists, then since G[Y] has minimum degree at least δ , we have that $G[X \setminus \{x_1, \ldots, x_{i-1}\}]$ has minimum degree at least $\delta - \delta/2 = \delta/2$. However, this contradicts part (2) of Lemma 3.8 since $|X \setminus \{x_1, \ldots, x_{i-1}\}| \leq |X| \leq \delta n/(1000d)$ and $\lambda \leq 10^{-4}\delta$.

Given this ordering we can do the following. Define an auxiliary bipartite graph H with parts A and B, where A consists of two copies of each vertex in X and B = Y. The edges are defined in the following manner - each copy of $x_i \in X$ in A has as neighbours those vertices in $\{x_1, \ldots, x_{i-1}\} \cup (Y \setminus X)$ which are neighbours of x_i in G. Crucially, it is easy to see that H has the following property.

Observation. If there is a matching in H covering A, then there is a forest \mathcal{F} in G[Y] which covers X and whose components are binary trees with leaves in $Y \setminus X$ and non-leaves in X.

Since our objective is to find such an \mathcal{F} , we need only to find a matching in H covering A - for this we will check that the conditions of Theorem 3.5 are satisfied. Consider any $S \subseteq A$ and let S' be the corresponding set in V(G) (not counting possible repetitions of a vertex in S) - note that $|S'| \geq |S|/2$. By the properties of the ordering x_1, x_2, \ldots, x_t of X we have that $e_G(N_H(S), S') \geq |S'|\delta/2 \geq |S|\delta/4$. Also, by part (1) of Lemma 3.8 we have $e_G(N_H(S), S') \leq \frac{d|N_H(S)||S'|}{n} + \lambda \sqrt{|S'||N_H(S)|} \leq \delta |N_H(S)|/500 + \lambda \sqrt{|S||N_H(S)|}$ since $|S'| \leq |S| \leq 2|X| \leq \delta n/500d$. Putting these together and using that $\delta \geq 100\lambda$ gives that $|N_H(S)| \geq |S|$ as desired.

Now that X has been absorbed into $Y \setminus X$, we can define the auxiliary random directed graph mentioned before. We define a directed graph H as follows. Denote the paths forming the linear forest \mathcal{F}' given by Lemma 6.4 by P_1, \ldots, P_l and the endpoints of each path P_i by $x_i, y_i \in Y \setminus X$ - let $U \subseteq Y \setminus X$ denote the set of all these endpoints. For each i, contract the pair $\{x_i, y_i\}$ into one vertex z_i and choose uniformly at random between defining $z_i^+ := x_i, z_i^- := y_i$ or $z_i^+ := y_i, z_i^- := x_i$. Let Z denote the set of vertices z_i - the vertex set of H is then $Z \cup Y \setminus (U \cup X)$. Note also that $|Z| \leq |X|$. The edges of H are defined in the following manner.

- 1. Every edge of $G[Y \setminus (U \cup X)]$ is oriented independently and uniformly at random.
- 2. For two vertices z_i, z_j , we include the directed edge $z_i \to z_j$ if $z_i^+ z_j^-$ is an edge of G.

3. For a vertex z_i and a vertex $v \in Y \setminus (U \cup X)$ we include the edge $z_i \to v$ (conversely, $v \to z_i$) if z_i^+v (conversely, z_i^-v) is an edge of G.

Crucially, by replacing each vertex z_i with the path P_i , it is easy to observe that H has the following deterministic property.

Observation 6.5. If there is a spanning linear forest in H, then there is a spanning linear forest \mathcal{F} of G[Y] with the same number of paths, all of which have their endpoints in $Y \setminus X$.

Therefore, we will need only to show that with positive probability, H satisfies the conditions of Lemma 6.3 with $r := \Theta(\delta)$. For convenience, we make first some parallel notation between H and G. For any set S consisting of vertices in V(H), we let S_G denote the set of vertices in $Y \subseteq V(G)$ given by substituting every $z_i \in S$ by the vertices x_i, y_i . Note that trivially $|S| \le |S_G| \le 2|S|$ and that for every sets $A, B \subseteq V(H)$ we have $\overrightarrow{e}_H(A, B) \le e_G(A_G, B_G)$.

We now want to show that H satisfies the conditions of Lemma 6.3. First, we note the following.

Lemma 6.6. With positive probability, the following all hold.

- 1. H has minimum out-degree and in-degree at least $\delta/10$.
- 2. For all $A, B \subseteq V(H)$ with $|A|, |B| \ge \delta n/100d$, there are at least $\delta |A|/4000$ edges directed from A to B.

Proof. We first deal with the degree conditions. Note that since by assumption, the set $Y \setminus X$ is δ -clean in G, it is easy to check that the definition of H implies that all vertices v of H are such that their outdegree and in-degree are random variables which stochastically dominate the binomial random variable $\text{Bin}(\delta, 1/2)$. Indeed, note that if $v \notin U$ has both elements of some pair $\{x_i, y_i\}$ as neighbours in G, then z_i will be both an in and out-neighbour of it in H; if only one of x_i, y_i is a neighbour, then with probability 1/2, it will be an in/out-neighbour in H. If $v = z_i$ for some i, then the analysis is similar. Therefore, the event E_v that v has out-degree or in-degree smaller than $\delta/10$ has $\mathbb{P}(E_v) = e^{-\Omega(\delta)}$. Further, since G is d-regular, the event E_v depends on at most $O(d^2)$ other events $E_{v'}$. Indeed it is easy to see that the events E_v is independent of all events $E_{v'}$ when v' is not a neighbour of v or there is no pair $\{x_i, y_i\}$ with both v, v' having neighbours in $\{x_i, y_i\}$. On the other hand when v, v' both have neighbours in the pair $\{x_i, y_i\}$, then, knowing the event $E_{v'}$ affects the choice of z_i^+, z_i^- and therefore also the out/in-degree of v. Because $e^{-\Omega(\delta)} = o(d^{-2})$, Lemma 3.4 now implies that with probability at least $\left(1 - \frac{1}{O(d^2)}\right)^n > e^{-n}$, no event E_v occurs and thus, the first condition holds.

For the second condition, we apply a union bound. Indeed, let $A, B \subseteq V(H)$ be sets of size at least $\delta n/100d$ - note that we have at most 4^n such pairs. Further, consider the sets $A' := A \setminus Z$ and $B' := B \setminus Z$, which, since $|Z| \leq |X| \leq \delta n/1000d$, have size at least $\delta n/200d$. The random procedure defining H implies that $\overrightarrow{e}_H(A', B')$ has the distribution of the binomial random variable Bin $(e_G(A', B'), 1/2)$. By Lemma 3.8 we have that

$$e_G(A', B') \ge d|A'||B'|/n - \lambda \sqrt{|A'||B'|} > d|A'||B'|/2n \ge \delta |A'|/400 > \delta |A|/1000 > 500n,$$

where the second inequality follows since $\delta \geq 10000\lambda$ and the last inequality since $\delta \geq 10000\lambda \geq 10000\sqrt{d/2}$ (recalling the remark after Theorem 3.7). So, by Lemma 3.3, with probability at most e^{-3n} we have that $\overrightarrow{e}_H(A', B') < \delta |A|/4000$. By a union bound, we then have that with probability at least $1 - 4^n \cdot e^{-3n} > 1 - e^{-n}$ all such pairs A, B have $\overrightarrow{e}_H(A, B) \geq \delta |A|/4000$.

Now, we can combine both conditions. Indeed, we showed that the first occurs with probability larger than e^{-n} and the second with probability larger than $1 - e^{-n}$. Thus, with positive probability both occur.

We can now prove Proposition 6.1.

Proof of Proposition 6.1. Assume that H satisfies the conditions 1 and 2 of the lemma above. Let us then verify the conditions of Lemma 6.3 with $r := \delta/10^{10}$. Take $H' \subseteq H$ to be a subdigraph such that $\Delta^+(H'), \Delta^-(H') < r$ and let $S \subseteq V(H), T := \Gamma^+_{H \setminus H'}(S)$. Suppose for sake of contradiction that |T| < |S|. First, note that by condition 1 of the lemma above, we have that $\overrightarrow{e}_H(S,T) \ge \delta|S|/20$ since the out-degree in $H \setminus H'$ is still at least $\delta/10 - r \ge \delta/20$. In turn Lemma 3.8 implies that

$$\overrightarrow{e}_{H}(S,T) \leq e_{G}(S_{G},T_{G}) \leq \frac{d|S_{G}||T_{G}|}{n} + \lambda \sqrt{|S_{G}||T_{G}|} < |S| \left(\frac{4d|S|}{n} + 2\lambda\right),$$

where we are using that $|S_G| \leq 2|S|$, $|T_G| \leq 2|T|$ and |S| > |T|. Observe now that the two inequalities above give a contradiction when $|S| \leq \delta n/100d$ since $\delta > 10000\lambda$. Hence, we can assume that $|S| > \delta n/100d$. Furthermore, note that the above argument also shows that if $|S| \leq \delta n/100d$, then $\Gamma_{H\backslash H'}^-(S)$ has size at least |S|, since condition 1 of the lemma above also gives that every vertex has in-degree at least $\delta/10$.

Now, suppose that $|S| > \delta n/100d$ and note that part 2 of Lemma 6.6 implies that $|T| > |V(H)| - \delta n/100d$. Indeed, we have for all sets B of size at least $\delta n/100d$ that $\overrightarrow{e}_H(S,B) > r|S|$ and so, $\overrightarrow{e}_{H\backslash H'}(S,B) \neq \emptyset$. Therefore, we can further assume that $|S| > |V(H)| - \delta n/100d$, since otherwise $|T| \geq |S|$. In this case since |T| < |S|, we can consider a set $T' \subseteq V(H) \setminus T$ of size $|V(H)| - |S| < |T'| \leq \delta n/100d$. By the previous paragraph, we then have that $\Gamma_{H\backslash H'}^-(T')$ has size at least |T'|. Also note that by definition of T', we have $\Gamma_{H\backslash H'}^-(T') \cap S = \emptyset$. This implies that $|T'| + |S| \leq |V(H)|$ which is a contradiction since we have |S| > |V(H)| - |T'| by assumption. Concluding, we must always have $|T| \geq |S|$ and so, Lemma 6.3 implies that H contains the desired spanning linear forest with positive probability. By Observation 6.5, this implies the existence of the desired linear forest \mathcal{F} in G[Y].

7 Hamilton cycles from good collections of paths

In this section, we will prove Theorem 1.3. More generally, we will prove that pseudorandom graphs which contain a certain spanning structure must also contain a Hamilton cycle. We will now describe this structure. First, given a path P and a collection of vertex-disjoint cycles $C = \{C_1, \ldots, C_l\}$, we say that P connects C if there exists a choice of edges $e_i \in C_i$ for every i so that P contains every edge e_i and is disjoint to the rest. Given such a setting, for each i we let $x_i^P, y_i^P \in C_i \setminus e_i$ denote the vertices of C_i which are not incident to e_i but are closest to this edge in the cycle - with the possibility that $x_i^P = y_i^P$ in the case that C_i only has three vertices. The set of these vertices is denoted by $V_{P,C} := \{x_i^P, y_i^P : 1 \le i \le l\}$ and we shall refer to it as the flexible set. As outlined in Section 2, the path P is able to absorb, for every i independently, the vertices of $V(C_i \setminus e)$ into the path, by replacing e_i with the path $C_i - e$. This flexibility allows us to connect partial structures outside of P. Since each connection requires logarithmic length in general, we are only able to do few connections using the flexible set, so we will assume in addition that the remaining vertices are already partitioned into few paths. This motivates the following definition.

Definition 7.1. A collection \mathcal{F} of vertex-disjoint paths in a graph G is said to be (r, l, δ) -good if it consists of r paths P_1, P_2, \ldots, P_r with the following properties.

- 1. The path P_1 connects a collection \mathcal{C} of l vertex-disjoint cycles, which are disjoint to the paths P_2, \ldots, P_r . Furthermore, these paths together with \mathcal{C} span the whole vertex set.
- 2. $V_{P_1,C}$ induces a graph with minimum degree at least δ .
- 3. The endpoints of the paths P_1, \ldots, P_r have each at least δ neighbours in the flexible set $V_{P_1,\mathcal{C}}$.

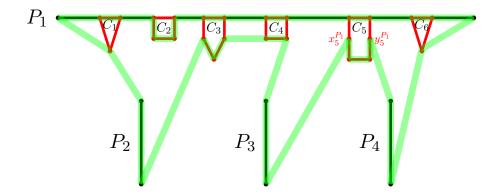


Figure 2: An illustration of a good collection of paths and how it can be used to form a Hamilton cycle. The path P_1 connects the cycles C_1, C_2, \ldots, C_6 and together with the paths P_2, P_3, P_4 it spans the whole vertex set. A Hamilton cycle is drawn in green by connecting the endpoints of the paths P_1, P_2, P_3, P_4 using the cycles. Those cycles which are not used for connecting (in this case C_2) are then absorbed into the path P_1 .

From the above definition and the lemmas presented in Section 3.5, we can show that pseudorandom graphs which contain a good collection of paths must also be Hamiltonian. The reader might want to refer to Figure 2 for an illustration.

Theorem 7.2. Let G be an (n, d, λ) -graph with $\lambda < d/500$. Let \mathcal{F} be an $\left(\frac{l}{100 \log n}, l, 500\lambda\right)$ -good collection of paths on V(G) for some $l \geq 2000\lambda n/d$. Then $G' := G \cup \mathcal{F}$ contains a Hamilton cycle.

Proof. Let \mathcal{F} consist of the paths P_1, \ldots, P_r and a cycle collection $\mathcal{C} = \{C_1, \ldots, C_l\}$. For each cycle C_j , let (x_j, y_j) denote the vertices $x_j^{P_1}, y_j^{P_1} \in C_j \setminus e_j$ as earlier defined and \mathcal{M} denote the collection of these pairs. Further, for each path P_i let u_i, v_i denote its endpoints and define the collection \mathcal{P} of pairs $\{a_j, b_j\}$ with $a_j := v_j, b_j := u_{j+1}$ (indices modulo r). Notice that we can apply Theorem 3.18. Indeed, $V(\mathcal{M}) = V_{P_1,\mathcal{C}}$, and so, the first and second conditions imply that $V(\mathcal{M})$ is 500λ -clean and $|\mathcal{M}| = l \geq 2000\lambda n/d$. Further, we have that $|\mathcal{P}| \leq \frac{|\mathcal{M}|}{100 \log n}$. Theorem 3.18 then implies that there exist vertex-disjoint \mathcal{M} -alternating paths between the pairs $\{a_j, b_j\}$ in the graph obtained by adding every non-existing edge $x_i y_i$ to G. In turn, it is easy to see that this produces a Hamilton cycle in G'.

7.1 Graphs with many cycles have a good collection of paths

As a corollary of Theorem 7.2, we can show that pseudorandom graphs which contain many cycles are Hamiltonian. Clearly, this will imply Theorem 1.3.

Theorem 7.3. Let G be an (n,d,λ) -graph with $d \ge (\log n)^{11}$ and $\lambda < d/10^{10}$. Suppose that it contains $20000\lambda n/d$ many vertex-disjoint cycles. Then it contains an $\left(\frac{l}{100\log n}, l, 500\lambda\right)$ -good collection of paths for some $l \ge 2000\lambda n/d$. Therefore, it is Hamiltonian.

Proof. Let G be an (n, d, λ) -graph and let $C = \{C_1, \ldots, C_t\}$ denote the given collection of t vertex-disjoint cycles, where $t \geq 20000\lambda n/d$. Notice that we can also assume that $t \geq n/(2\log n)$ since such a cycle collection will always exist if say, $d \geq 100\lambda$ (this is easy to see by greedily finding cycles of size at most $\log n$). We will now first prove a series of preparatory lemmas in order to make this collection more suitable to work with. The following ensures that every vertex has many neighbours disjoint to these.

Lemma 7.4. There exists a subcollection $C' \subseteq C$ of at least 3t/4 cycles such that every vertex in G has at least d/10 neighbours not in V(C').

Proof. First, we refine \mathcal{C} so that it only contains short cycles. Indeed, at most t/100 of the cycles have length larger than 100n/t. Let us delete all of these from \mathcal{C} . Let now $\mathcal{C}' \subseteq \mathcal{C}$ be randomly chosen by letting every cycle be included in it independently and with probability 4/5. Using Lemma 3.3, we then have that with high probability, $|\mathcal{C}'| \geq 3t/4$ as required. Further, every vertex has degree $d \geq (\log n)^{11}$ and since each cycle is now of length at most $100n/t = O(\log n)$ (since by before, we can assume that $t \geq n/(2\log n)$) we can apply Lemma 3.3 with k := 100n/t, to get that the probability that a vertex has less than d/10 neighbours not in $V(\mathcal{C}')$ is $o(n^{-1})$. By a union bound over all vertices, we then have also that with high probability, every vertex has at least d/10 neighbours not in $V(\mathcal{C}')$. Combining both of these properties gives the desired outcome for some choice of \mathcal{C}' .

Given the above, let us redefine \mathcal{C} to be \mathcal{C}' . We now need the following claim which deals with finding a path connecting a decent proportion of the cycles in \mathcal{C} .

Lemma 7.5. There exists a subcollection $C' \subseteq C$ of t/2 cycles and a path P connecting them. Furthermore, among every two consecutive edges of P there is an edge belonging to some cycle in C'.

Proof. For each cycle $C_i \in \mathcal{C}$, pick an arbitrary edge $e_i = x_i y_i$ belonging to it. Let \mathcal{M} denote the collection of pairs $\{x_i, y_i\}$. Now, by part (4) of Lemma 3.8, every two disjoint subsets of G of size $\lambda n/d$ have an edge between them. Therefore, Lemma 3.17 implies that $G \cup \mathcal{M}$ contains an \mathcal{M} -alternating path which uses all but at most $2\lambda n/d - 1$ of the edges e_i . Since $|\mathcal{C}| - 2\lambda n/d \ge 3t/4 - 2\lambda n/d \ge t/2$, this gives the desired path P.

Observe now that since among every two consecutive edges of P there is an edge belonging to some cycle of C', we have from Lemma 7.4 that every vertex has at least d/10 neighbours not in $V(P) \cup V(C')$. In order to continue let us again redefine C to be C'.

Now, let us denote \mathcal{C} by $\{C_1,\ldots,C_r\}$ with $r\geq t/2$ and that the edges e_1,\ldots,e_r appear in this order in the path P. We further delete the cycles $C_1,\ldots,C_{r/3},C_{2r/3},\ldots,C_r$ from \mathcal{C} so that all the cycles now in consideration are present in the middle of the path P. More precisely, no cycle in \mathcal{C} will intersect the first $r/3\geq 2\lambda n/d$ or last $r/3\geq 2\lambda n/d$ vertices of the path P - let us mark this as property (*). Now, since $r/3\geq t/6\geq 3000\lambda n/d$, Lemma 3.12 directly implies the following.

Lemma 7.6. There exists a subcollection $C' \subseteq C$ of at least t/7 cycles such that $V_{P,C'}$ is 500λ -clean.

Again, let us redefine C to be the new collection C', which is the final collection of cycles - it has size now $l := |C| \ge t/7 \ge 2000 \lambda n/d$ and $|V_{P,C}| \ge l$ as well. Let us note that by Lemma 3.11, at most $4\lambda^2 n^2/d^2|V_{P,C}| \le \lambda n/d$ vertices have degree less than $d|V_{P,C}|/2n \ge dl/2n \ge 500\lambda$ in $V_{P,C}$. Therefore, among the first $2\lambda n/d$ vertices of P there is a vertex x with degree at least 500λ in $V_{P,C}$ and among the last vertices of P there is also such a vertex y. By property (*), we can then make P an xy-path and still have that P connects C, so that its endpoints have degree at least 500λ in $V_{P,C}$.

Finally, having constructed $P_1 := P$ and the collection \mathcal{C} , we are left to find the paths P_2, P_3, \ldots, P_r and ensure the third condition of Definition 7.1. For this we will apply Proposition 6.1. Note first that $Y := V(G) \setminus (V(P) \cup V(\mathcal{C}))$ is, by before, such that G[Y] has minimum degree at least d/10 (which by Lemma 3.8 clearly implies that $|Y| \geq n/11$). Moreover, let $X' := \{v \in Y : d(v, V_{P,\mathcal{C}}) < 500\lambda\}$. As previously noted, we have that $|X'| \leq \lambda n/d$. Furthermore, the second part of Lemma 3.11 applied to $S := Y \setminus X'$ implies that there is $X' \subseteq X \subseteq Y$ with $|X| \leq |X'| + \lambda n/d \leq 2\lambda n/d$ such that $Y \setminus X$ is clean. Since as noted before, $|Y| \geq n/11$, then $G[Y \setminus X]$ has minimum degree at least d/100. Applying Proposition 6.1 to X, Y with $\delta := d/100$ implies the existence of the desired paths P_2, \ldots, P_r .

To finish, note that this proposition gives us, together with P_1 and C, a $\left(O\left(n/\delta^{1/5}\right), l, 500\lambda\right)$ -good collection of paths. Since $t \geq n/2 \log n$, $l \geq t/7 \geq n/14 \log n$ and $d \geq (\log n)^{11}$, we have $O\left(n/\delta^{1/5}\right) \leq \frac{l}{100 \log n}$, as desired.

As stated before, Theorem 1.3 is an easy corollary of the above.

Proof of Theorem 1.3. Let G be an (n,d,λ) -graph with $\lambda < \alpha d/10^{10}$, $d \ge n^{\alpha}$ and $n \ge 10000$. We claim that G contains at least $\alpha n/8$ vertex-disjoint cycles of size at most $4/\alpha$. Indeed, we can do this in a greedy manner - suppose that we have vertex-disjoint cycles C_1, \ldots, C_r of size at most $4/\alpha$. Then, if $r < \alpha n/8$, the set X of vertices not contained in any of these cycles has size at least n/2. By part (2) of Lemma 3.8 it must then be that G[X] has at least $\frac{d|X|^2}{2n} - \frac{\lambda n}{2} \ge \frac{nd}{10} > n^{1+\alpha/2}$ edges. By Theorem 3.6, G[X] contains a cycle C_{r+1} of length at most $4/\alpha$ and this can be added to the collection of vertex-disjoint cycles.

We can then apply Theorem 7.3 to G since $\alpha n/8 \geq 20000 \lambda n/d$, thus showing that G is Hamiltonian. \square

We finish this section by noting that clearly Theorem 7.3 has the slightly more general corollary than Theorem 1.3 concerning different densities of our pseudorandom graph. Its proof is the same as that of Theorem 1.3.

Corollary 7.7. Let G be an (n,d,λ) -graph with $d \ge (\log n)^{11}$ and $d/\lambda > 10^{10} \log_d n$. Then, G is Hamiltonian.

Indeed, note that in general, it is easy to see that greedily we can always find a collection of $n/4 \log_d n$ vertex-disjoint cycles. Therefore, if $d/\lambda > 10^{10} \log_d n$, then the conditions of Theorem 7.3 are satisfied.

8 Applications

In this section, we discuss in more detail the applications of our new techniques and main results, that we mentioned briefly in Section 1.1.

8.1 Random Cayley graphs and Hamilton cycles with few colours

Let G be a group and $A \subseteq G$ a subset. The Cayley graph $\Gamma(G,A)$ is the graph with vertex set G and all edges of the form $\{g,ga\}$ for $g \in G, a \in A$. Note that we consider here undirected Cayley graphs, hence, $\Gamma(G,A)$ is a d-regular graph with $|A| \leq d = |A \cup A^{-1}| \leq 2|A|$. We will assume here that A does not contain the neutral element, so that $\Gamma(G,A)$ has no loops. As discussed in Section 1.1, studying the Hamiltonicity of Cayley graphs (Conjecture 1.5) is a very important partial case of the Lovász conjecture mentioned in Section 1. Up to now, only special cases of Conjecture 1.5 have been proven, most notably when the ambient group is Abelian or when A is linear in the size of G (see Christofides, Hladký and Máthé [16]). But the general case is still wide open. In view of this, a lot of research has been invested into understanding random Cayley graphs, where G is a group of order n and $A \subseteq G$ of size d = d(n) is chosen uniformly at random, and the question is whether $\Gamma(G,A)$ is Hamiltonian with high probability (as $n \to \infty$).

Alon and Roichman [6] proved that for $d = O(\log n)$, whp $\Gamma(G, A)$ is connected (in fact, they proved the much stronger statement that it is an expander). This bound is tight up to a constant factor, for instance, in $G = \mathbb{Z}_2^n$, the smallest size of a set A such that $\Gamma(G, A)$ is connected is $\log_2 |G|$. The analogous question for Hamiltonicity has turned out to be much more difficult. As a first result, Meng and Huang [42] showed that almost all Cayley graphs are Hamiltonian. By using the technique of Alon and Roichman to bound the second eigenvalue of a random Cayley graph, Krivelevich and Sudakov [33] deduced from their bound on the Hamiltonicity of (n, d, λ) -graphs that already $O(\log^5 n)$ random generators suffice for Hamiltonicity. Their technique for applying the Alon and Roichman result was later refined by Christofides and Markström [15] to show that $O(\log^3 n)$ random generators suffice. It was also noted that Conjecture 1.1 would imply that $O(\log n)$ generators suffice, which seems reasonable since the random Cayley graph is then connected as discussed above, and there is no obvious obstacle to Hamiltonicity. This was also stated as a conjecture by Pak and Radoičić [43].

Conjecture 8.1. Let C be a sufficiently large absolute constant. Let G be a group of order n and $d = C \log n$. If $A \subseteq G$ is a set of size d chosen uniformly at random, then with high probability, $\Gamma(G, A)$ is Hamiltonian.

Using an operator Hoeffding inequality together with our new Theorem 1.2, we can improve the number of random generators to $O(\log^{5/3} n)$.

Theorem 8.2. Let C be a sufficiently large absolute constant. Let G be a group of order n and $d = C \log^{5/3} n$. If $A \subseteq G$ is a set of size d chosen uniformly at random, then with high probability, $\Gamma(G, A)$ is Hamiltonian.

A closely related problem concerns the quest for finding Hamilton cycles with few colours in optimally edge-coloured complete graphs. This is a proper edge-colouring of K_n which uses n-1 colours when n is even and n colours when n is odd. The following conjecture was proposed by Akbari, Etesami, Mahini, and Mahmoody [2].

Conjecture 8.3. Every properly edge-coloured K_n with $\chi'(K_n)$ colours has a Hamilton cycle with $O(\log n)$ colours.

Akbari et al. proved that one can find a Hamilton cycle with at most $8\sqrt{n}$ colours. The bound on the number of colours was later improved to $O(\log^3 n)$ by Balla, Pokrovskiy and Sudakov [9]. Their strategy was to pick $O(\log^3 n)$ colours at random and show that the subgraph induced by these colours is already Hamiltonian with high probability. Here, we obtain a further improvement to $O(\log^{5/3} n)$.

Theorem 8.4. Every properly edge-coloured K_n with $\chi'(K_n)$ colours has a Hamilton cycle with $O(\log^{5/3} n)$ colours.

We prove both Theorems 8.2 and 8.4 in a unified way. Note that when n is even, then an optimal edgecolouring of K_n is a partition of the edge set into perfect matchings. Moreover, if G is a group and $a \in G$ is not the neutral element, then the edges in the Cayley graph corresponding to a and a^{-1} form either a 1-factor (if $a = a^{-1}$) or a 2-factor (if $a \neq a^{-1}$). It turns out that whenever K_n is partitioned into spanning regular graphs and we choose sufficiently many of them randomly, then with high probability, the obtained graph is pseudorandom and hence Hamiltonian.

Theorem 8.5. For all R > 0 there exists C > 0 such that the following holds: Assume K_n is edge-partitioned into regular spanning subgraphs H_1, \ldots, H_t , with degrees $1 \le r_i \le R$ for all $i \in [t]$. Let c_1, \ldots, c_k be a sequence of indices chosen independently from [t] such that i is always chosen with probability at least 1/Rt. If $k \ge C(\log n)^{5/3}$, then with high probability, $G = \bigcup_{j \in [k]} H_{c_j}$ is Hamiltonian.

Proof. Since we do not require the c_j 's to be uniformly distributed, we first reduce to this case by introducing a dummy variable * and a straightforward coupling. For each $j \in [k]$ independently, we define c_j' to be a random element from $[t] \cup \{*\}$ as follows. If $c_j = i$ then $c_j' = i$ with probability $\frac{1/\mathbb{P}[c_j = i]}{Rt}$ and $c_j' = *$ with probability $1 - \frac{1/\mathbb{P}[c_j = i]}{Rt}$. Clearly this is well defined since $\mathbb{P}[c_j = i] \ge 1/Rt$ for all i. Thus we have that $\mathbb{P}[c_j' = i] = 1/Rt$ for all $i \in [t]$ and $\mathbb{P}[c_j' = *] = 1 - 1/R$. Let $A^{(j)}$ be the adjacency matrix of H_{c_j} with probability $\frac{n-1}{n}$ and $A^{(j)} = r_{c_j}I$ with probability 1/n. Moreover, let Y_j be the 0-matrix if $c_j' = *$ and $Y_j = \frac{1}{4R}(A^{(j)} - \frac{r_{c_j'}}{n}J)$ otherwise, where J is the all 1's matrix. Crucially, we have $\mathbb{E}[Y_j] = 0$. Indeed, the sum of the adjacency matrices of all H_i is J - I and c_j' is uniformly distributed over [t] if we condition on $c_j' \neq *$. On the other hand, if $c_j' = *$ then $Y_j = 0$. Moreover, note that Y_j is a symmetric matrix with all eigenvalues in [-1/2, 1/2]. Indeed, note that $A^{(j)}$ has eigenvalues in [-R, R] since H_{c_j} is either the adjacency matrix of a graph with maximum degree at most R or $r_{c_j}I$; thus, $\frac{1}{4R}A^{(j)}$ has eigenvalues in [-1/4, 1/4]. Also, clearly $\frac{1}{4R} \cdot \frac{r_{c_j'}}{n}J$ has eigenvalues in [-1/4, 1/4].

Using this setup for each individual $j \in [k]$, we now let $I \subseteq [k]$ be the set of j for which $c'_j = c_j$. Obviously, $G' = \bigcup_{j \in I} H_{c_j}$ is a subgraph of G, so it suffices to establish Hamiltonicity of G'. We set further $A = \sum_{j \in I} A^{(j)}$ and $d = \sum_{j \in I} r_{c_j}$. Observe that $X_i = Y_1 + \cdots + Y_i$ is a martingale and $X_k = \frac{1}{4R}(A - \frac{d}{n}J)$. Using an operator Hoeffding inequality for Hilbert spaces, developed by Christofides and Markström [14] (see [9, Theorem 4] for the version we use), one can show that for any $0 < \varepsilon < 1/2$, it holds that

$$\mathbb{P}[||X_k|| \ge \varepsilon k] \le 2n \exp(-2\varepsilon^2 k). \tag{3}$$

We may assume that $k = \lceil C(\log n)^{5/3} \rceil$. Plugging $\varepsilon = \sqrt{\frac{\log n}{k}} \le \frac{1}{\sqrt{C}(\log n)^{1/3}}$ into (3), we see that with high probability,

 $||A - \frac{d}{n}J|| = 4R||X_k|| \le 4\varepsilon Rk.$

Moreover, with high probability, the c_j 's with $j \in I$ are pairwise distinct and $A^{(j)} \neq r_{c_j}I$. Finally, by a simple Chernoff bound, we have $|I| \geq k/2R$ with high probability.

We assume now that all these events hold. Then, crucially, A is the adjacency matrix of the random graph G' and $\lambda(G') = ||A - \frac{d}{n}J|| \le 4\varepsilon Rk$. So we can conclude that G' is an (n,d,λ) -graph with $d = |I| \ge k/2R$ and $d/\lambda \ge 1/8\varepsilon R^2$. Since $\varepsilon \le \frac{1}{\sqrt{C}(\log n)^{1/3}}$, by choosing C large enough, we can apply our Theorem 1.2 to establish Hamiltonicity of G'.

We can now easily deduce the two results stated above.

Proof of Theorem 8.2. Let G be a group of order n. Instead of choosing A of order d uniformly at random, we may choose d elements $a_1, \ldots, a_d \in G$ independently and uniformly at random. For each $a \in G$ which is not the neutral element, let H_a be the spanning subgraph of the complete graph on G with all edges of the form $\{g, ga\}$ for $g \in G$. As discussed earlier, $H_a = H_{a^{-1}}$ is either a 1-factor (if $a = a^{-1}$) or a 2-factor (if $a \neq a^{-1}$). Each such factor is chosen with probability $\Theta(1/n)$ when we pick a_i , hence Theorem 8.5 implies the claim.

Proof of Theorem 8.4. Balla, Pokrovskiy and Sudakov [9] found a reduction for odd n to the case when n is even. Thus, it suffices to consider the case when n is even. Then, an optimal edge-colouring of K_n is simply a partition into perfect matchings. We choose $O(\log^{5/3} n)$ of these at random. With high probability, the graph obtained is Hamiltonian, in particular, there is a Hamilton cycle with only $O(\log^{5/3} n)$ colours.

8.2 Additive patterns in multiplicative subgroups

As we mentioned in Section 1.1, we can also apply our results to a problem of Alon and Bourgain [4] on additive patterns in multiplicative subgroups. It is well-known that any multiplicative subgroup A of the finite field \mathbb{F}_q of size at least $q^{3/4}$ must contain two elements x, y such that x + y also belongs to A. Alon and Bourgain [4] studied more complex additive structures in multiplicative subgroups and proved that when A as above has size even $|A| \ge q^{3/4}(\log q)^{1/2-o(1)}$, then there is a cyclic ordering of the elements of A such that the sum of any two consecutive elements is also in A. We improve on this result by showing that the additional polylog-factor can be avoided.

Theorem 8.6. There exists an absolute positive constant c so that for any prime power q and for any multiplicative subgroup A of the finite field \mathbb{F}_q of size $|A| \geq cq^{3/4}$ there is a cyclic ordering of the elements of A such that the sum of any two consecutive elements is also in A.

The proof is essentially the same as that of Alon and Bourgain [4, Theorem 1.2], we only need to apply our new Theorem 1.3 instead of the previous Krivelevich and Sudakov bound.

Proof. Let c be a sufficiently large constant. Let G be the graph with vertex set \mathbb{F}_q where $xy \in E(G)$ if and only if $x+y \in A$ (this is the so-called Cayley sum graph). By definition, G has q vertices, is |A|-regular and it was shown in [4] (see Lemma 2.7) that its second in absolute value eigenvalue is at most $q^{1/2}$. Now, let H = G[A]. It is easy to check that H is regular as well. Let d denote the degree of the vertices in H. Applying Lemma 3.8 to G, we see that $e_G(A) \geq |A|^3/2q - q^{1/2}|A|$, so by choosing c large enough we have $d \geq |A|^2/2q$. Moreover, by interlacing of eigenvalues, it follows that H is a $(|A|, d, q^{1/2})$ -graph. Since $d \geq \frac{1}{2}c^2q^{1/2}$, we can then apply Theorem 1.3 to obtain a Hamilton cycle in H, which yields the desired structure.

8.3 Classification-free proof of a result of Pak and Radoicic

As mentioned in Section 1.1, Pak and Radoičić [43] proved that every finite group G of size at least 3 has a generating set S of size $|S| \leq \log_2 |G|$ such that the corresponding Cayley graph $\Gamma(G, S)$ contains a Hamilton cycle. However, their proof relies on the Classification of Finite Simple Groups, and motivated by the results discussed in Section 8.1, they asked if there exists a classification-free proof of their result. We essentially prove this, though our result is not quite as sharp since we can only guarantee a set S of size $O(\log |G|)$. For the proof, we use a combination of a random and a deterministic choice of generators. As in Section 8.1, we will use the fact that $O(\log |G|)$ random generators form with high probability a pseudorandom graph. The decisive advantage which we have here is to deterministically pick, say, three generators which will give us linearly many vertex-disjoint triangles. Together with an additional argument we can then facilitate a direct application of Theorem 7.2.

Theorem 8.7. Every group G of order n has a generating set S of size $O(\log n)$ such that $\Gamma(G,S)$ is Hamiltonian.

Proof. Let C be a sufficiently large constant. We let S consist of three parts S_1, S_2, S_3 . First, let $a, b \in G$ be distinct with $b \neq a^{-1}$, and set $S_1 = \{a, b, (ab)^{-1}\}$. We claim that $\Gamma(G, S_1)$ contains n/10 vertex-disjoint triangles. Indeed, suppose that C_1, \ldots, C_r are vertex-disjoint triangles and let $X \subseteq G$ be the set of vertices covered by those. If r < n/10, then there is $g \in G$ such that $g, ga, gab \notin X$, which gives us an additional triangle that we can add to the collection. Let from now on t := n/10 denote the number of triangles.

Next, let $S_2 \subseteq G$ be a random set of size $d = C \log n$. Proceeding analogously to the proof of Theorem 8.5, with high probability, $\Gamma(G, S_2)$ contains an (n, d, λ) -graph H with $\lambda \leq \varepsilon d$ for $\varepsilon := 10^{-10}$. In particular, such a choice exists, which we will fix from now on.

We will now proceed similarly as in the proof of Theorem 7.3. Analogously to Lemmas 7.4–7.6, we can find a path P which connects a subcollection C of at least $t/7 \ge n/70$ many vertex-disjoint triangles found above with the following properties:

- $V_{P,C}$ is clean in H, has size $t/7 \ge n/70$ and thus spans a subgraph with minimum degree at least $d|V_{P,C}|/2n \ge d/200$;
- the endpoints of P have degree at least d/200 into $V_{P,C}$;
- every vertex has at least d/10 neighbours in Y, where Y is the complement of $V(P) \cup V(\mathcal{C})$.

It remains to cover Y with a linear forest with few paths all of whose endpoints have many neighbours in $V_{P,C}$. We will first deal with the vertices that are not suitable endpoints. Let X be the set of vertices which have fewer than $d/200 \le d|V_{P,C}|/2n$ neighbours in $V_{P,C}$. By Lemma 3.11, $|X| \le 4\lambda^2 n^2/d^2|V_{P,C}| \le n/10^{17}$. Hence, we can apply Lemma 6.4 with $\delta := d/200$ to obtain a linear forest \mathcal{F}_0 which covers X such that all its paths have their endpoints outside X. We let \mathcal{F} denote the union of this linear forest together with the path P and regarding all remaining vertices as trivial paths of length 0. This means we have almost found the desired good collection of paths. The only problem is that the number of paths in this

linear forest is potentially too big (for example, it can be linear in n). To this end, we will now *combine* the paths of \mathcal{F} to reduce their number to at most $n/10000 \log n$. By *combining*, we simply mean that we can add an edge between the endpoints of two paths to replace them with a longer path. Crucically, this preserves the property that the endpoints of the paths are "good" in the sense that they still have large degree into $V_{P,\mathcal{C}}$.

Let \mathcal{F} be the current linear forest. We claim that there exists a set S_3 of $O(\log n)$ generators such that using the edges of $\Gamma(G, S_3)$, we can combine paths of \mathcal{F} to obtain a linear forest with at most $n/10000 \log n$ paths. The idea is to find S_3 iteratively by always adding a generator which allows to combine as many paths as possible. We will show that this strategy performs as desired by analysing it in dyadic steps. To this end, we claim that if we have found a linear forest with at most n/k paths, for some k, then we can choose at most 10k additional generators to reduce the number of paths to at most n/2k. Indeed, as long as the number of paths is at least n/2k, give each path a direction and let A be the set of all starting vertices and B the set of all terminal vertices (if a path consists only of one vertex, we add a copy of that vertex to both sets). Now, consider the bipartite graph with classes A, B where two vertices are adjacent if and only if they are ends of distinct paths. Let $m=|A|=|B|\geq n/2k$. This graph has m^2-m edges, each corresponding to an edge in the complete graph on G. By averaging, there exists a generator a such that $\Gamma(G,\{a\})$ contains at least $(m^2-m)/n$ edges from this auxiliary graph. We pick such a generator. Note that in this set of edges, we can pick at least a third of all edges to form a matching of size at least $n/10k^2$. Adding these matching edges to our collection of paths, we find a new linear forest and the number of paths has been reduced by at least $n/10k^2$. Repeating this procedure for at most 10k steps will yield a linear forest with at most n/2k paths, as desired.

Now, we can simply apply the above claim iteratively with $k = 1, 2, 4, ..., 2^j$ to see that a linear forest with at most $n/2^j$ paths has been found when the set S_3 has size at most $10 \left(1 + 2 + 4 + \cdots + 2^j\right) \leq 20 \cdot 2^j$. Finally, we can apply Theorem 7.2 to conclude that, setting $S := S_1 \cup S_2 \cup S_3$, the Cayley graph $\Gamma(G, S)$ is Hamiltonian.

9 Concluding remarks

In this paper we developed several new techniques and results concerning the Hamiltonicity of pseudorandom graphs. The main open problem here is to prove that every (n, d, λ) -graph with $d/\lambda \geq C$, for some universal constant C > 0, has a Hamilton cycle. We proved this in the case that d is a small polynomial of n. The key idea was to find linearly many vertex-disjoint cycles that we could arrange into an absorbing structure. While this proof method does not immediately extend to the sparser setting, we believe that many of the techniques we developed could play a crucial role in future work concerning Conjecture 1.1. In particular, the use of the Friedman-Pippenger tree embedding technique with rollbacks and our method to cover a small set of 'bad' vertices with paths having their endpoints outside the bad set seem very useful for a potential proof of Conjecture 1.1 in full.

In the general case, we were able to reduce the required ratio d/λ to order $\log^{1/3} n$. In particular, our results can be applied to pseudo-random graphs with sub-logarithmic degree. One very important example of such graphs are random graphs. As we discussed in the introduction, the Hamiltonicity of random graphs is well understood. Particularly, the celebrated result of Pósa [44] states that the random graph G(n, p) is Hamiltonian with high probability when $p \geq C \log n/n$. In addition to the rotation-extension technique, his proof relies on so-called 'booster' edges, which is where the randomness of the graph is crucially used. Our arguments in the proof of Theorem 1.2 can in fact be slightly adapted to prove the following statement, which gives another proof of Pósa's result.

Theorem 9.1. Let G be an n-vertex graph satisfying part (1) of Lemma 3.8 with $\lambda = \frac{d}{C(\log n)^{1/3}}$ for some large constant C and such that every vertex has degree between d/100 and 100d. Then, G is Hamiltonian.

This provides a proof of Pósa's result which uses no direct randomness. Indeed, it is not difficult to check that G(n, p) for $p \ge 100 \log n/n$ satisfies the conditions of Theorem 9.1 with high probability.

Finally, we note that the case of presudo-random graphs with constant degree d remains widely open. In fact, even when G is (n,d,λ) -graph which is almost optimally pseudorandom (i.e., $\lambda = O(\sqrt{d})$), it is not known whether it contains a nearly spanning path, i.e., a path of length n - o(n). It is not too difficult to show that G contains a path of length $(1 - \varepsilon)n$ for some small constant $\varepsilon = \varepsilon_d$. Indeed, one can use Pósa's rotation-extension technique to show that (n,d,λ) -graphs always contain paths of length at least $\left(1 - \frac{100\lambda^2}{d^2}\right)n$ provided that $d > 10\lambda$. However, it is not clear how this can be improved and solving this problem is a natural step towards proving Conjecture 1.1.

Acknowledgments. The authors would like to thank Nemanja Draganić for fruitful discussions regarding the Friedman-Pippenger technique.

References

- [1] M. Ajtai, J. Komlós, and E. Szemerédi. First occurrence of Hamilton cycles in random graphs. *North-Holland Math. Stud.*, 115(C):173–178, 1985.
- [2] S. Akbari, O. Etesami, H. Mahini, and M. Mahmoody. On rainbow cycles in edge colored complete graphs. Australas. J. Combin., 37:33, 2007.
- [3] N. Alon. The linear arboricity of graphs. Israel J. Math., 62(3):311–325, 1988.
- [4] N. Alon and J. Bourgain. Additive patterns in multiplicative subgroups. Geom. Funct. Anal., 24:721–739, 2014.
- [5] N. Alon and F. R. K. Chung. Explicit construction of linear sized tolerant networks. *Discrete Math.*, 72(1-3): 15–19, 1988.
- [6] N. Alon and Y. Roichman. Random Cayley graphs and expanders. Random Structures Algorithms, 5(2): 271–284, 1994.
- [7] N. Alon and J. H. Spencer. The probabilistic method. John Wiley & Sons, 2016.
- [8] L. Babai. Long cycles in vertex-transitive graphs. J. Graph Theory, 3(3):301–304, 1979.
- [9] I. Balla, A. Pokrovskiy, and B. Sudakov. A remark on Hamilton cycles with few colors. *Moscow Journal of Combinatorics and Number Theory*, 7(3):73–77, 2017.
- [10] I. Ben-Eliezer, M. Krivelevich, and B. Sudakov. The size Ramsey number of a directed path. *J. Combin. Theory Ser. B*, 102(3):743–755, 2012.
- [11] B. Bollobás. Modern graph theory. Springer Science & Business Media, 184, 1998.
- [12] B. Bollobás. The evolution of sparse graphs. Graph theory and combinatorics (Cambridge, 1983), 35–57, 1984.
- [13] B. Bollobás, T. I. Fenner, and A. M. Frieze. An algorithm for finding Hamilton paths and cycles in random graphs. *Combinatorica*, 7:327–341, 1987.
- [14] D. Christofides and K. Markström. Expansion properties of random Cayley graphs and vertex transitive graphs via matrix martingales. *Random Structures Algorithms*, 32(1):88–100, 2008.
- [15] D. Christofides and K. Markstrom. Random Latin square graphs. Random Structures Algorithms, 41(1):47–65, 2012.
- [16] D. Christofides, J. Hladký, and A. Máthé. Hamilton cycles in dense vertex-transitive graphs. J. Combin. Theory Ser. B, 109:34–72, 2014.
- [17] V. Chvátal and P. Erdős. A note on Hamiltonian circuits. Discrete Math., 2(2):111–113, 1972.
- [18] C. Cooper, A. Frieze, and B. Reed. Random regular graphs of non-constant degree: connectivity and Hamiltonicity. *Combin. Probab. Comput.*, 11(3):249–261, 2002.
- [19] B. Csaba, D. Kühn, A. Lo, D. Osthus, and A. Treglown. Proof of the 1-factorization and Hamilton decomposition conjectures. *Mem. Amer. Math. Soc.*, 244, monograph 1154, 164 pages, 2016.
- [20] B. Cuckler and J. Kahn. Hamiltonian cycles in Dirac graphs. Combinatorica, 29:299–326, 2009.
- [21] S. J. Curran and J. A. Gallian. Hamiltonian cycles and paths in Cayley graphs and digraphs—a survey. *Discrete Math.*, 156(1-3):1–18, 1996.
- [22] M. DeVos. Longer cycles in vertex transitive graphs. arXiv:2302.04255, 2023.

- [23] G. A. Dirac. Some theorems on abstract graphs. Proc. Lond. Math. Soc., 3(1):69-81, 1952.
- [24] N. Draganić, M. Krivelevich, and R. Nenadov. Rolling backwards can move you forward: on embedding problems in sparse expanders. *Trans. Amer. Math. Soc.*, 375(07):5195–5216, 2022.
- [25] A. Ferber, E. Long, and B. Sudakov. Counting Hamilton decompositions of oriented graphs. Int. Math. Res. Not. IMRN, (22):6908–6933, 2018.
- [26] J. Friedman and N. Pippenger. Expanding graphs contain all small trees. Combinatorica, 7:71–76, 1987.
- [27] A. Frieze and M. Krivelevich. Hamilton cycles in random subgraphs of pseudo-random graphs. *Discrete Math.*, 256(1-2):137–150, 2002.
- [28] R. J. Gould. Recent advances on the Hamiltonian problem: Survey III. Graphs Combin., 30:1–46, 2014.
- [29] D. Hefetz, M. Krivelevich and T. Szabó. Hamilton cycles in highly connected and expanding graphs. Combinatorica, 29(5):547–568, 2009.
- [30] F. Knox, D. Kühn, and D. Osthus. Approximate Hamilton decompositions of random graphs. *Random Structures Algorithms*, 40(2):133–149, 2012.
- [31] J. Komlós and E. Szemerédi. Limit distribution for the existence of Hamiltonian cycles in a random graph. Discrete Math., 43(1):55–63, 1983.
- [32] A. Korshunov. Solution of a problem of Erdős and Rényi on Hamilton cycles in non-oriented graphs. *Soviet Math Dokl*, 17:760–764, 1976.
- [33] M. Krivelevich and B. Sudakov. Sparse pseudo-random graphs are Hamiltonian. *J. Graph Theory*, 42(1): 17–33, 2003.
- [34] M. Krivelevich. The critical bias for the Hamiltonicity game is (1+o(1)). J. Amer. Math. Soc., 24(1):125–131, 2011.
- [35] M. Krivelevich and B. Sudakov. Pseudo-random graphs. In *More sets, graphs and numbers*, pages 199–262. Springer, 2006.
- [36] M. Krivelevich, B. Sudakov, V. H. Vu, and N. C. Wormald. Random regular graphs of high degree. *Random Structures Algorithms*, 18(4):346–363, 2001.
- [37] M. Krivelevich, C. Lee, and B. Sudakov. Robust Hamiltonicity of Dirac graphs. *Trans. Amer. Math. Soc.*, 366 (6):3095–3130, 2014.
- [38] D. Kühn and D. Osthus. Hamilton decompositions of regular expanders: a proof of Kelly's conjecture for large tournaments. *Adv. Math.*, 237:62–146, 2013.
- [39] D. Kühn and D. Osthus. Hamilton cycles in graphs and hypergraphs: an extremal perspective. In Proceedings of the International Congress of Mathematicians 2014, Seoul, Korea., 4:381–406, 2014. pages 381–406. Kyung Moon Sa, Seoul, 2014.
- [40] D. Kühn, J. Lapinskas, D. Osthus, and V. Patel. Proof of a conjecture of Thomassen on Hamilton cycles in highly connected tournaments. *Proc. Lond. Math. Soc.*, 109(3):733–762, 2014.
- [41] L. Lovász. Combinatorial structures and their applications. In *Proc. Calgary Internat. Conf.*, Calgary, Alberta, 243–246, 1970.
- [42] J. Meng and Q. Huang. Almost all Cayley graphs are Hamiltonian. Acta Math. Sinica, 12(2):151–155, 1996.
- [43] I. Pak and R. Radoičić. Hamiltonian paths in Cayley graphs. Discrete Math., 309(17):5501-5508, 2009.
- [44] L. Pósa. Hamiltonian circuits in random graphs. Discrete Math., 14(4):359–364, 1976.
- [45] E. Rapaport Strasser. Cayley color groups and Hamilton lines. Scripta Math., 24:51–58, 1959.